

Assignment 2 - Scientific Computing SC9505

Jared Wogan

2022-11-21

All code was compiled inside an Ubuntu WSL environment on an AMD 3900X with 32GB of RAM. All timing and performance data was run with dense, randomly generated arrays and vectors. A full repository containing all of my code can be found at <https://github.com/JaredWogan/SC9505>.

1 Matrix Vector Multiplication Using MPI

This program calculates the product of a matrix with a vector. Our task is to determine the benefits (if any) that parallelization provides for this problem. The parallelized code can be confirmed to be working, as seen in Figure 1.

The program has been parallelized as follows. First, the main process will initialize both the matrix and the vector that will be used in the calculations. The main process then broadcasts the vector to all other threads. In order to perform the full calculation, the main process will send each thread a unique row of the matrix, which the thread will then use to calculate the corresponding element of the result vector. These results are collected by the main process, which will subsequently send the threads a new row, if there is more work to be done.

```
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MV$ make clean
rm -rf 15-matrix-vector.o 16-serial.o
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MV$ make all
mpiCC -O2 -o 15-matrix-vector.o 10-matrix-vector.cpp
mpiCC -O2 -o 16-serial.o 11-serial-matrix-vector.cpp
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MV$ mpirun -n 4 ./15-matrix-vector.o 1 10 10
Average calculation time = 0.0000002603
Total time = 0.0001655470
( 0, 1, 2, 3, 4, 5, 6, 7, 8, 9)
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MV$ mpirun -n 8 ./15-matrix-vector.o 1 20 20
Average calculation time = 0.0000002964
Total time = 0.0003846000
( 0, 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16, 17, 18, 19)
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MV$
```

Figure 1: Demonstration of the code working. Here the matrix is taken to be the identity, and the vector is filled with consecutive integers.

The scaling when the program is run on 4 threads as a function of the number of rows/columns was determined to be $\mathcal{O}(N^{1.5})$, see Figure 2. This is less than is expected for matrix vector multiplication,

which is typically an $\mathcal{O}(N^2)$ operation. The serial code I used to compare the parallel code with scaled roughly as $\mathcal{O}(N^{2.2})$.

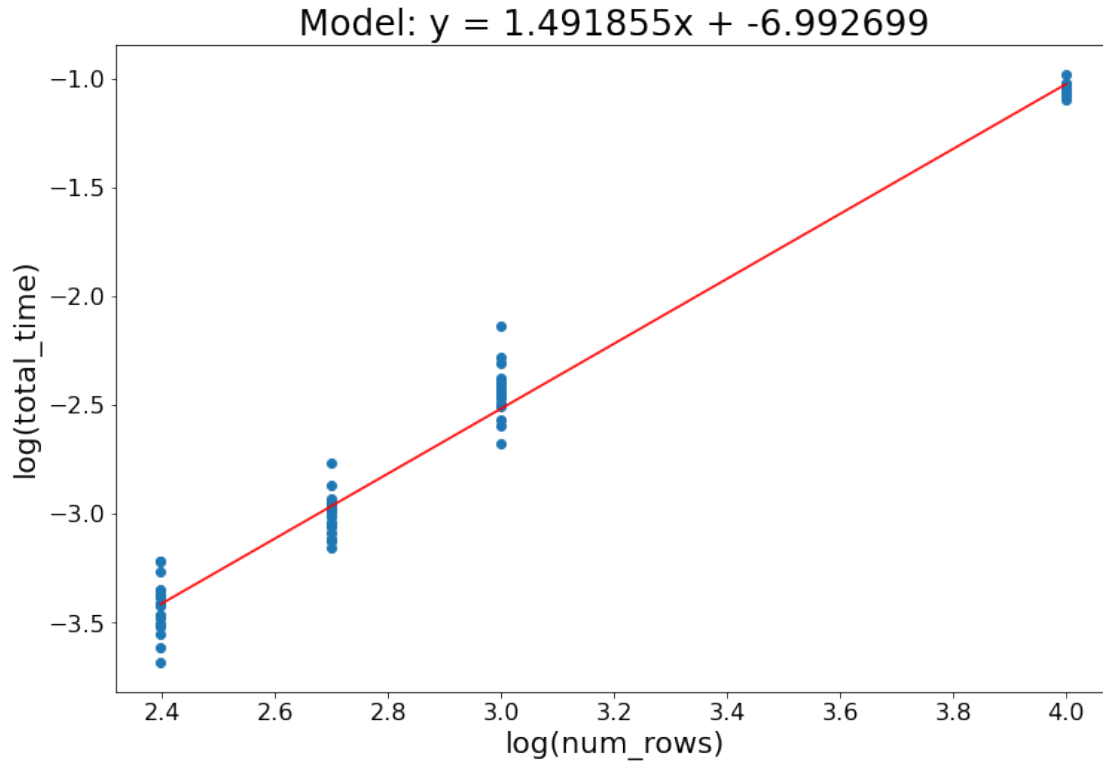


Figure 2: Plotting the runtime of the program as a function of the number of rows N (on a log-log scale), we find that the total runtime scales as $\mathcal{O}(N^{1.5})$, which is less than the expected $\mathcal{O}(N^2)$. Note the plot is generated with data for a constant amount of threads, which in this case is 4.

Looking at the scaling instead for a constant number of rows (here we are using 10000), we find there is no real scaling. Ideally we would like to see the total runtime decrease as we increase the number of threads, but this is not the case for the algorithm being used here.

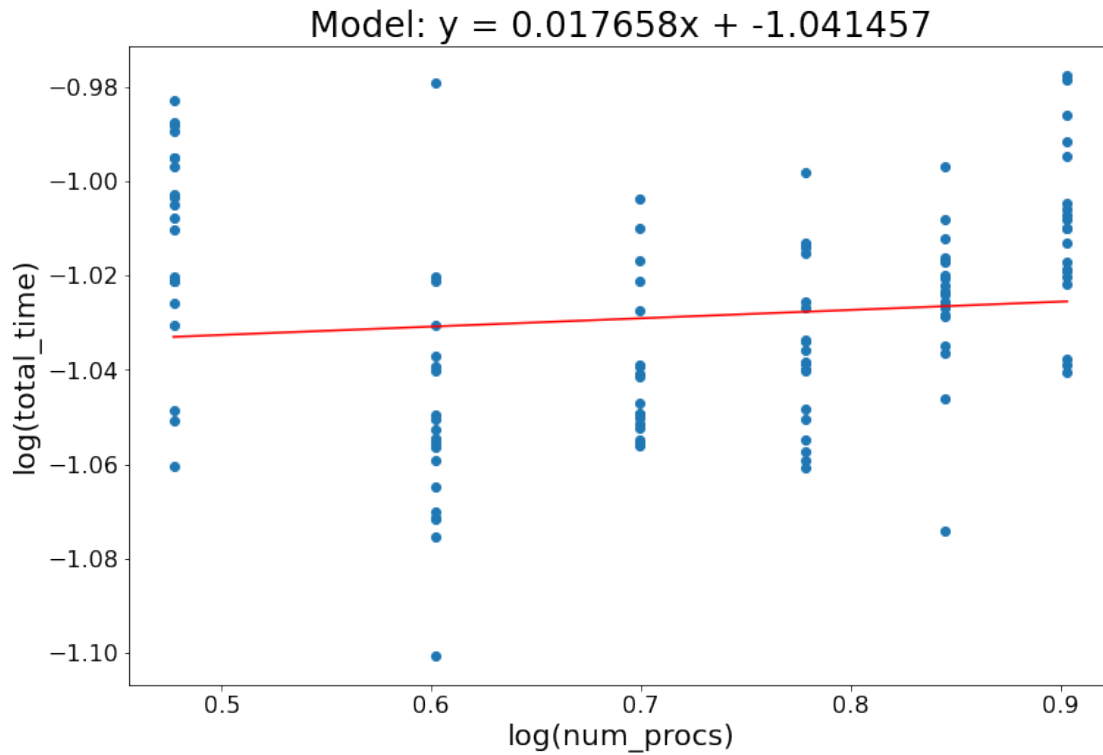


Figure 3: Plotting the runtime of the program as a function of the number of threads P (on a log-log scale). The linear fit is not ideal, such that it is fair to conclude that there is no note worthy scaling with the number of processors. Again it is worth noting that the plot has been generated for a constant number of rows which was taken to be 10000.

As discussed in lecture, the efficiency of our algorithm is less than ideal, which is supported by the calculations in Table 1 (see Appendix A for details on how these values were calculated). We see that running the code across 3 threads yields superlinear scaling, and our program is more efficient than the serial code. However, as the number of threads increase, we start seeing diminishing returns in terms of the speedup, and so our efficiency starts dropping dramatically. At 8 concurrent threads, our code is less than half as efficient as the serial code, and it is actually slower than the case with only 3 threads (in reality they are approximately equal in runtime, but the linear model predicted it would be slower).

If we wanted to write a program to do matrix vector multiplication in parallel more efficiently, it would be best to have a shared pool of memory that each process can read and write to at the same time. This could be done by adding a single line to the serial code and using OpenMP.

Number of Threads	Speedup	Scaling Factor a	Efficiency
3	3.92	1.31	1.31
4	3.76	0.94	0.94
5	3.61	0.72	0.72
6	3.47	0.58	0.58
7	3.33	0.48	0.48
8	3.19	0.40	0.40

Table 1: Summary of the speedup and efficiency of our matrix vector multiplication code. These values are taken for a constant number of rows $N = 10000$.

```

1  //////////////////////////////////////
2  // Matrix-vector multiplication code Ab=c //
3  //////////////////////////////////////
4
5  // Note that I will index arrays from 0 to n-1.
6  // Here workers do all the work and boss just handles collating results
7  // and sending info about A.
8
9  // include, definitions, globals etc here
10 #include <iostream>
11 #include <fstream>
12 #include <iomanip>
13 #include <random>
14 #include "boost/multi_array.hpp"
15 #include "mpi.h"
16
17 using namespace std;
18
19 class MPI_Obj
20 {
21 public:
22     int size;
23     int rank;
24
25     MPI_Obj(int &argc, char **&argv)
26     {
27         MPI_Init(&argc, &argv);
28         MPI_Comm_size(MPI_COMM_WORLD, &size);
29         MPI_Comm_rank(MPI_COMM_WORLD, &rank);
30     }
31
32     ~MPI_Obj()
33     {
34         MPI_Finalize();
35     }
36 };
37
38 void GetArraySize(int &output, int &nrows, int &ncols, MPI_Obj &the_mpi, int argc, char **argv)
39 {
40     // Accept input parameters from the command line, otherwise ask the user
41     if (the_mpi.rank == 0)
42     {
43         if (argc == 4)
44         {
45             output = atoi(argv[1]);
46             nrows = atoi(argv[2]);
47             ncols = atoi(argv[3]);
48         }
49         else
50         {
51             cout << "Please enter the number of rows -> ";
52             cin >> nrows;
53             cout << "Please enter the number of columns -> ";
54             cin >> ncols;
55         }
56     }
57
58     // send everyone nrows, ncols
59     int buf[2] = {nrows, ncols};
60     MPI_Bcast(buf, 2, MPI_INT, 0, MPI_COMM_WORLD);
61     if (the_mpi.rank != 0)

```

```

62     {
63         nrows = buf[0];
64         ncols = buf[1];
65     }
66 }
67
68 void SetupArrays(int nrows, int ncols, boost::multi_array<double, 2> &A, vector<double> &b,
69 ↪ vector<double> &c, vector<double> &Arow, MPI_Obj &the_mpi)
70 {
71     uniform_real_distribution<double> unif(-1.0, 10);
72     default_random_engine re;
73     // Boss part
74     if (the_mpi.rank == 0)
75     {
76         // Set size of A
77         A.resize(boost::extents[nrows][ncols]);
78
79         // Initialize A
80         for (int i = 0; i < nrows; ++i)
81             for (int j = 0; j < ncols; ++j)
82             {
83                 // Identity
84                 // if (i == j)
85                 //     A[i][j] = 1.0;
86                 // else
87                 //     A[i][j] = 0.0;
88                 // Reverse Identity
89                 // if (i == (ncols - j - 1))
90                 //     A[i][j] = 1.0;
91                 // else
92                 //     A[i][j] = 0.0;
93                 // Random
94                 A[i][j] = unif(re);
95             }
96
97         // Initialize b
98         for (int i = 0; i < ncols; ++i)
99         {
100             // b[i] = 1.0;
101             // b[i] = (double) i;
102             b[i] = unif(re);
103         }
104
105         // Allocate space for c, the answer
106         c.reserve(nrows);
107         c.resize(nrows);
108     }
109     // Worker part
110     else
111     {
112         // Allocate space for 1 row of A
113         Arow.reserve(ncols);
114         Arow.resize(ncols);
115     }
116
117     // send b to every worker process, note b is a vector so b and &b[0] not same
118     MPI_Bcast(&b[0], ncols, MPI_DOUBLE, 0, MPI_COMM_WORLD);
119 }
120
121 void Output(int output, vector<double> &c, MPI_Obj &the_mpi)
122 {
123     if (the_mpi.rank == 0 && output == 1)
124     {
125         cout << "( " << c[0];
126         for (int i = 1; i < c.size(); ++i)
127             cout << ", " << c[i];
128         cout << ")\n";
129     }
130 }
131
132 int main(int argc, char **argv)
133 {
134     // Data File
135     // ofstream datafile("/home/jared/Desktop/mv-timings.txt", ios_base::app);
136     ofstream datafile("mv-timings.txt", ios_base::app);
137
138     // initialize MPI

```

```

138 MPI_Obj the_mpi(argc, argv);
139 if (the_mpi.size < 2)
140     MPI_Abort(MPI_COMM_WORLD, 1);
141
142 // determine/distribute size of arrays here
143 int output = 1, nrows = 0, ncols = 0;
144 GetArraySize(output, nrows, ncols, the_mpi, argc, argv);
145 if (the_mpi.size - 1 > nrows)
146 {
147     MPI_Abort(MPI_COMM_WORLD, -1);
148 }
149
150 boost::multi_array<double, 2> A;
151 vector<double> b(ncols);
152 vector<double> c;
153 vector<double> Arow;
154 SetupArrays(nrows, ncols, A, b, c, Arow, the_mpi);
155
156 MPI_Status status;
157
158 // Timing variables
159 double calc_time = 0, avg_time, total_time;
160
161 // Boss part
162 if (the_mpi.rank == 0)
163 {
164     total_time = MPI_Wtime();
165     // send one row to each worker tagged with row number, assume size<nrows
166     int rowsent = 0;
167     for (int i = 1; i < the_mpi.size; i++)
168     {
169         MPI_Send(&A[rowsent][0], ncols, MPI_DOUBLE, i, rowsent + 1, MPI_COMM_WORLD);
170         rowsent++;
171     }
172
173     for (int i = 0; i < nrows; i++)
174     {
175         double ans;
176         MPI_Recv(&ans, 1, MPI_DOUBLE, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
177         int sender = status.MPI_SOURCE;
178         int anstype = status.MPI_TAG; // row number+1
179         c[anstype - 1] = ans;
180         if (rowsent < nrows)
181         { // send new row
182             MPI_Send(&A[rowsent][0], ncols, MPI_DOUBLE, sender, rowsent + 1, MPI_COMM_WORLD);
183             rowsent++;
184         }
185         else
186         { // tell sender no more work to do via a 0 TAG
187             MPI_Send(MPI_BOTTOM, 0, MPI_DOUBLE, sender, 0, MPI_COMM_WORLD);
188         }
189     }
190 }
191 else
192 {
193     // Get a row of A
194     MPI_Recv(&Arow[0], ncols, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
195     while (status.MPI_TAG != 0)
196     {
197         // work out Arow.b
198         double ans = 0.0;
199
200         calc_time -= MPI_Wtime();
201         for (int i = 0; i < ncols; i++)
202             ans += Arow[i] * b[i];
203         calc_time += MPI_Wtime();
204
205         // Send answer of Arow.b back to boss and get another row to work on
206         MPI_Send(&ans, 1, MPI_DOUBLE, 0, status.MPI_TAG, MPI_COMM_WORLD);
207         MPI_Recv(&Arow[0], ncols, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
208     }
209 }
210
211 if (the_mpi.rank == 0)
212     total_time = MPI_Wtime() - total_time;
213
214 MPI_Reduce(&calc_time, &avg_time, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);

```

```

215
216     if (the_mpi.rank == 0)
217     {
218         avg_time /= (the_mpi.size - 1); // Boss node doesn't do any of the calculations
219         printf("Average calculation time = %.10f\n", avg_time);
220         printf("Total time = %.10f\n", total_time);
221
222         datafile << fixed << setprecision(10);
223         datafile << nrows << " " << ncols << " " << the_mpi.size << " " << avg_time << " " << total_time
↪ << endl;
224     }
225
226     // output c here on Boss node
227     Output(output, c, the_mpi);
228 }

```

2 Matrix Matrix Multiplication Using MPI

Our second task is to generalize the matrix vector product code to perform matrix matrix products. The modified code can be seen below, while a demonstration that the code works correctly can be seen in Figure 4. The serial code that will be used for comparison was written using the DGEMM_ function from the BLAS library.

The program that computes a matrix matrix product has been parallelized by first having the main process assemble the two matrices to be used. Each thread is subsequently sent a copy of the second matrix by the main process. To compute the product, each thread is sent a row of the first matrix, which it will then use to compute an entire column of the resulting matrix. Upon completion, the thread will send the completed column to the main process in exchange for a new row (if there are any left).

```
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MM$ make all
mpiCC -O2 -o 15-matrix-matrix.o 10-matrix-matrix.cpp
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MM$ mpirun -n 4 ./15-matrix-matrix.o 1 10 10 10
Average calculation time = 0.000029873
Total time = 0.000256180

1.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 2.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 3.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 4.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 5.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 6.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 7.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 8.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 9.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 10.0000

(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MM$ make all
mpiCC -O2 -o 15-matrix-matrix.o 10-matrix-matrix.cpp
(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MM$ mpirun -n 4 ./15-matrix-matrix.o 1 10 10 10
Average calculation time = 0.000022110
Total time = 0.000339140

0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 10.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 9.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 8.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
1.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000

(base) jared@Nebulae:~/Scientific Computing SC9505/Assignment 2 - November 30/MM$
```

Figure 4: Demonstration of the code working. In the first example, the first matrix is taken to be the identity and the second has consecutive integers along the main diagonal. In the second example, the first matrix is changed to the reverse identity (1s along the antidiagonal). In both cases, we get the expected result.

The modified code for matrix matrix multiplication scales as $\mathcal{O}(N^{3.3})$, where N is the number of rows as can be seen in Figure 5. This is slightly more than the expected $\mathcal{O}(N^3)$. It should be noted that the serial code scaled as $\mathcal{O}(N^{2.9})$, which is slightly less than expected. It is also interesting that the current world record (as of 2020) for the asymptotic scaling of matrix matrix multiplication has been shown to be of the order $\mathcal{O}(N^{2.3728596})$ by Josh Alman and Virginia Vassilevska Williams (arXiv:2010.05846).

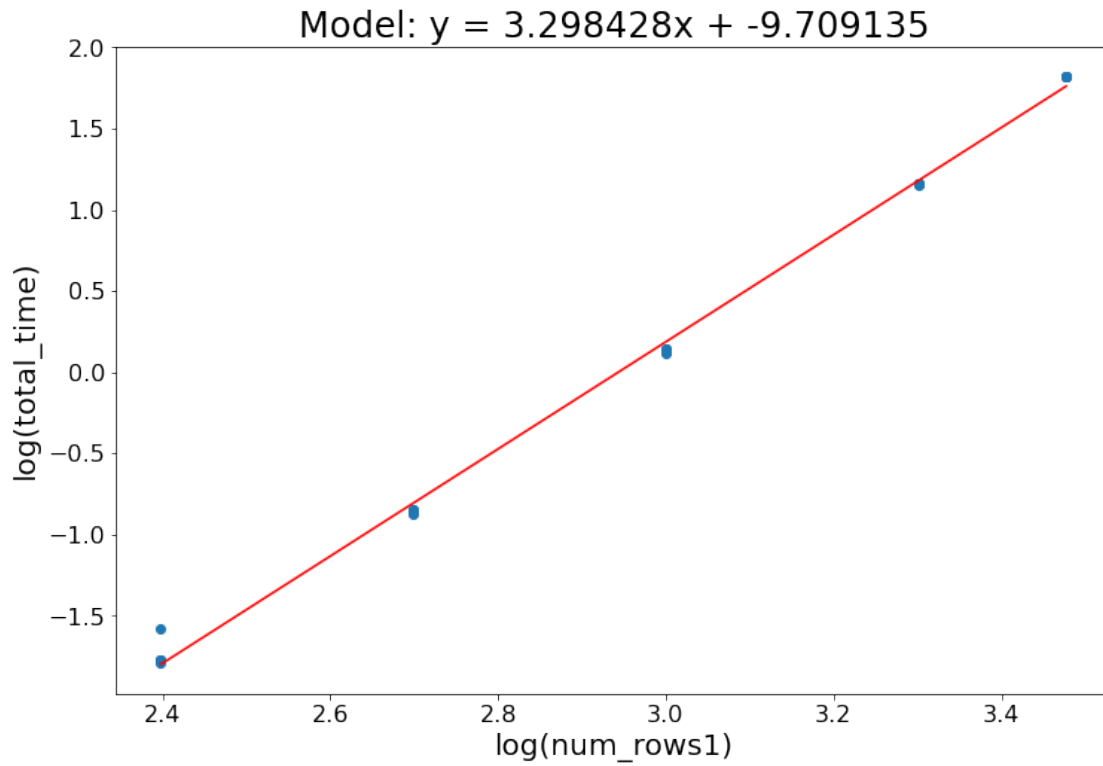


Figure 5: Plotting the runtime of the program as a function of the number of rows N (on a log-log scale), we find that the total runtime scales as $\mathcal{O}(N^{3.3})$, which is slightly more than the expected $\mathcal{O}(N^3)$.

More importantly for our purposes, we find that the total runtime scales as $\mathcal{O}(P^{-1})$, where P is the total number of threads (for a fixed number of rows N). This is the expected result as we discussed in lecture: for large N , we expect the speedup to approach P , which is supported by Figure 6 and Table 2.

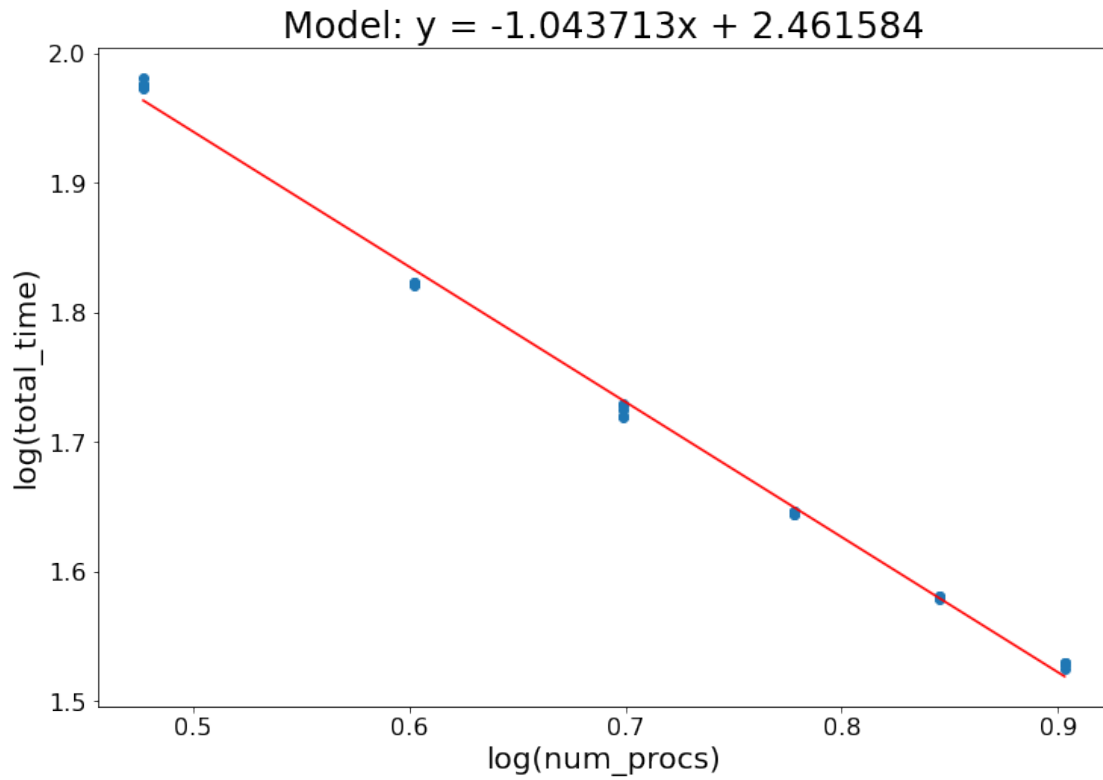


Figure 6: Plotting the runtime of the program as a function of the number of threads P , we find the scaling for a fixed number of rows N (here $N = 3000$) to be of the order $\mathcal{O}(P^{-1})$

Number of Threads	Speedup	Scaling Factor a	Efficiency
3	1.66	0.55	0.55
4	2.24	0.56	0.56
5	2.83	0.57	0.57
6	3.42	0.57	0.57
7	4.02	0.57	0.57
8	4.62	0.58	0.58

Table 2: Summary of the speedup and efficiency of our matrix matrix multiplication code. These values are taken for a constant number of rows $N = 3000$.

```

1  //////////////////////////////////////
2  // Matrix-vector multiplication code Ab=c //
3  //////////////////////////////////////
4
5  // Note that I will index arrays from 0 to n-1.
6  // Here workers do all the work and boss just handles collating results
7  // and sending info about A.
8
9  // include, definitions, globals etc here
10 #include <iostream>
11 #include <fstream>
12 #include <iomanip>
13 #include <random>
14 #include "boost/multi_array.hpp"
15 #include "mpi.h"
16

```

```

17 using namespace std;
18
19 class MPI_Obj
20 {
21 public:
22     int size;
23     int rank;
24
25     MPI_Obj(int &argc, char **&argv)
26     {
27         MPI_Init(&argc, &argv);
28         MPI_Comm_size(MPI_COMM_WORLD, &size);
29         MPI_Comm_rank(MPI_COMM_WORLD, &rank);
30     }
31
32     ~MPI_Obj()
33     {
34         MPI_Finalize();
35     }
36 };
37
38 void GetArraySize(int &output, int &nrows, int &ncols, MPI_Obj &the_mpi, int argc, char **argv)
39 {
40     if (the_mpi.rank == 0)
41     {
42         if (argc == 4)
43         {
44             output = atoi(argv[1]);
45             nrows = atoi(argv[2]);
46             ncols = atoi(argv[3]);
47         }
48         else
49         {
50             cout << "Please enter the number of rows -> ";
51             cin >> nrows;
52             cout << "Please enter the number of columns -> ";
53             cin >> ncols;
54         }
55     }
56
57     // send everyone nrows, ncols
58     int buf[2] = {nrows, ncols};
59     MPI_Bcast(buf, 2, MPI_INT, 0, MPI_COMM_WORLD);
60     if (the_mpi.rank != 0)
61     {
62         nrows = buf[0];
63         ncols = buf[1];
64     }
65 }
66
67 void SetupArrays(int nrows, int ncols, boost::multi_array<double, 2> &A, vector<double> &b,
68 ↪ vector<double> &c, vector<double> &Arow, MPI_Obj &the_mpi)
69 {
70     uniform_real_distribution<double> unif(-1.0, 10);
71     default_random_engine re;
72     // Boss part
73     if (the_mpi.rank == 0)
74     {
75         // Set size of A
76         A.resize(boost::extents[nrows][ncols]);
77
78         // Initialize A
79         for (int i = 0; i < nrows; ++i)
80             for (int j = 0; j < ncols; ++j)
81             {
82                 // Identity
83                 if (i == j)
84                     A[i][j] = 1.0;
85                 else
86                     A[i][j] = 0.0;
87                 // Reverse Identity
88                 // if (i == (ncols - j - 1))
89                 //     A[i][j] = 1.0;
90                 // else
91                 //     A[i][j] = 0.0;
92                 // Random
93                 // A[i][j] = unif(re);

```

```

93         }
94
95         // Initialize b
96         for (int i = 0; i < ncols; ++i)
97         {
98             // b[i] = 1.0;
99             b[i] = (double) i;
100             // b[i] = unif(re);
101         }
102
103         // Allocate space for c, the answer
104         c.reserve(nrows);
105         c.resize(nrows);
106     }
107     // Worker part
108     else
109     {
110         // Allocate space for 1 row of A
111         Arow.reserve(ncols);
112         Arow.resize(ncols);
113     }
114
115     // send b to every worker process, note b is a vector so b and &b[0] not same
116     MPI_Bcast(&b[0], ncols, MPI_DOUBLE, 0, MPI_COMM_WORLD);
117 }
118
119 void Output(int output, vector<double> &c, MPI_Obj &the_mpi)
120 {
121     if (the_mpi.rank == 0 && output == 1)
122     {
123         cout << "( " << c[0];
124         for (int i = 1; i < c.size(); ++i)
125             cout << ", " << c[i];
126         cout << ")\n";
127     }
128 }
129
130 int main(int argc, char **argv)
131 {
132     // Data File
133     // ofstream datafile("/home/jared/Desktop/mv-timings.txt", ios_base::app);
134     ofstream datafile("mv-timings.txt", ios_base::app);
135
136     // initialize MPI
137     MPI_Obj the_mpi(argc, argv);
138     if (the_mpi.size < 2)
139         MPI_Abort(MPI_COMM_WORLD, 1);
140
141     // determine/distribute size of arrays here
142     int output = 1, nrows = 0, ncols = 0;
143     GetArraySize(output, nrows, ncols, the_mpi, argc, argv);
144     if (the_mpi.size - 1 > nrows)
145     {
146         MPI_Abort(MPI_COMM_WORLD, -1);
147     }
148
149     boost::multi_array<double, 2> A;
150     vector<double> b(ncols);
151     vector<double> c;
152     vector<double> Arow;
153     SetupArrays(nrows, ncols, A, b, c, Arow, the_mpi);
154
155     MPI_Status status;
156
157     // Timing variables
158     double calc_time = 0, avg_time, total_time;
159
160     // Boss part
161     if (the_mpi.rank == 0)
162     {
163         total_time = MPI_Wtime();
164         // send one row to each worker tagged with row number, assume size<nrows
165         int rowsent = 1;
166         for (int i = 1; i < the_mpi.size; i++)
167         {
168             MPI_Send(&A[rowsent - 1][0], ncols, MPI_DOUBLE, i, rowsent, MPI_COMM_WORLD);
169             rowsent++;

```

```

170     }
171
172     for (int i = 0; i < nrows; i++)
173     {
174         double ans;
175         MPI_Recv(&ans, 1, MPI_DOUBLE, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
176         int sender = status.MPI_SOURCE;
177         int row = status.MPI_TAG - 1;
178         c[row] = ans;
179         if (rowsent - 1 < nrows)
180         { // send new row
181             MPI_Send(&A[rowsent - 1][0], ncols, MPI_DOUBLE, sender, rowsent, MPI_COMM_WORLD);
182             rowsent++;
183         }
184         else
185         { // tell sender no more work to do via a 0 TAG
186             MPI_Send(MPI_BOTTOM, 0, MPI_DOUBLE, sender, 0, MPI_COMM_WORLD);
187         }
188     }
189 }
190 else
191 {
192     // Get a row of A
193     MPI_Recv(&Arow[0], ncols, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
194     while (status.MPI_TAG != 0)
195     {
196         // work out Arow.b
197         double ans = 0.0;
198
199         calc_time -= MPI_Wtime();
200         for (int i = 0; i < ncols; i++)
201             ans += Arow[i] * b[i];
202         calc_time += MPI_Wtime();
203
204         // Send answer of Arow.b back to boss and get another row to work on
205         MPI_Send(&ans, 1, MPI_DOUBLE, 0, status.MPI_TAG, MPI_COMM_WORLD);
206         MPI_Recv(&Arow[0], ncols, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
207     }
208 }
209
210 if (the_mpi.rank == 0)
211     total_time = MPI_Wtime() - total_time;
212
213 MPI_Reduce(&calc_time, &avg_time, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
214
215 if (the_mpi.rank == 0)
216 {
217     avg_time /= (the_mpi.size - 1); // Boss node doesn't do any of the calculations
218     printf("Average calculation time = %.10f\n", avg_time);
219     printf("Total time = %.10f\n", total_time);
220
221     datafile << fixed << setprecision(10);
222     datafile << nrows << " " << ncols << " " << the_mpi.size << " " << avg_time << " " << total_time
223     << endl;
224 }
225
226 // output c here on Boss node
227 Output(output, c, the_mpi);

```

3 Matrix Matrix Multiplication Using MPI and BLAS

In this section, we modify the code to compute the matrix matrix product from above. The new code is parallelized identically to the non-BLAS code, with the only change being the method each thread uses locally to compute the product of a row with the second matrix. Namely, we use the `DDOT_` function provided by the BLAS library. As we did for the previous code, we test the code still works as can be seen in Figure 7.

```
(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$ make all
rm -rf 15-matrix-matrix-blas.o
(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$ make all
mpiCC -O2 -o 15-matrix-matrix-blas.o 10-matrix-matrix-blas.cpp -lblas -llapack
(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$ mpirun -n 4 ./15-matrix-matrix-blas.o 1 10 10 10
Average calculation time = 0.0000359777
Total time = 0.0000887870

1.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 2.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 3.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 4.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 5.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 6.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 7.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 8.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 9.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 10.0000

(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$ make all
mpiCC -O2 -o 15-matrix-matrix-blas.o 10-matrix-matrix-blas.cpp -lblas -llapack
(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$ mpirun -n 4 ./15-matrix-matrix-blas.o 1 10 10 10
Average calculation time = 0.0000048020
Total time = 0.0000310880

0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 10.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 8.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 7.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 6.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 0.0000, 5.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 0.0000, 4.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 0.0000, 3.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
0.0000, 2.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000
1.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000, 0.0000

(base) jared@NebruLae:~/Scientific Computing SC9505/Assignment 2 - November 30/MMblas$
```

Figure 7: Demonstration of the code working. Just as we did for the non-BLAS implementation, the first example takes the first matrix to be the identity while the second has consecutive integers along the main diagonal. In the second example, the first matrix is changed to the reverse identity (1s along the antidiagonal). In both cases, we get the expected result.

The new code using BLAS shows that the total runtime as a function of the number of rows N scales as $\mathcal{O}(N^{3.6})$ as supported by Figure 8. This is peculiar since this means the modified code using BLAS actually scales worse than the non-BLAS code.

Perhaps this is because the BLAS code doesn't actually scale in polynomial time and is instead scaling like $\mathcal{O}(N^\alpha \log N)$ for some α , but I am not confident enough to say for certain¹. I would expect the BLAS code to scale better than the non-BLAS code, and given that the linear regression does not fit as nicely as it did for the non-BLAS case, I suspect this may be the case.

¹If there is a more obvious or apparent reason for this behaviour, I'd be interested to know.

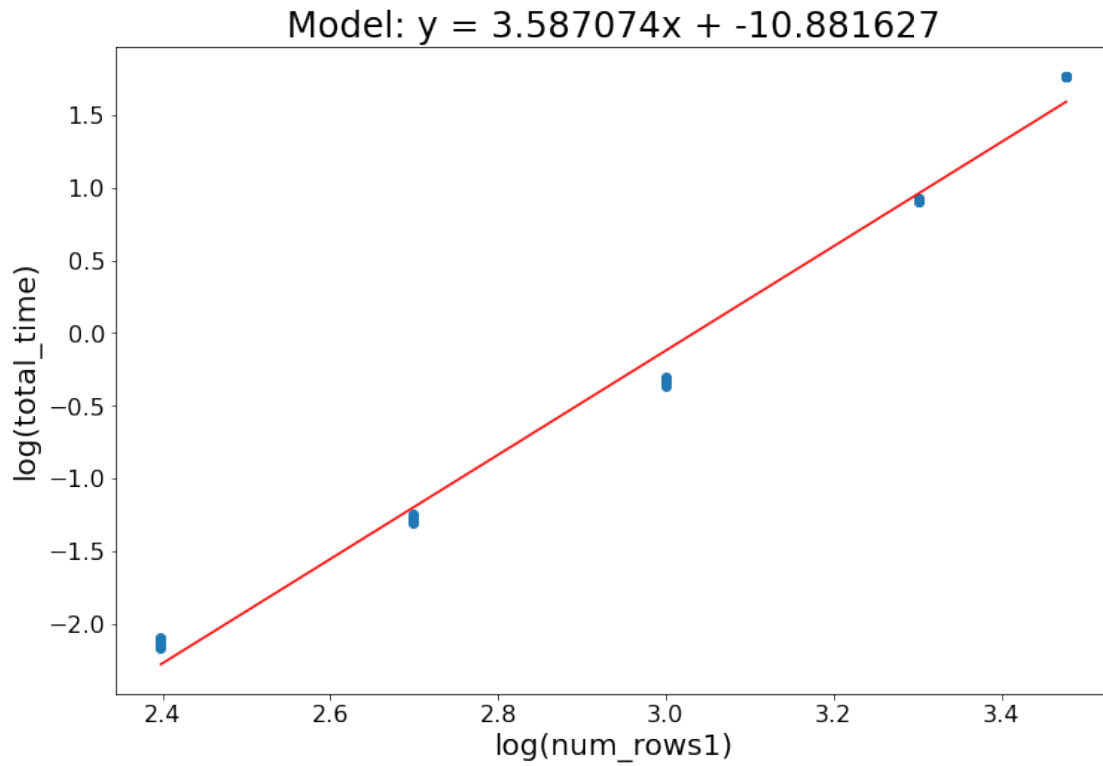


Figure 8: Plotting the runtime of the program as a function of the number of rows N (on a log-log scale), we find that the total runtime scales as $\mathcal{O}(N^{3.6})$, which again is more than the expected $\mathcal{O}(N^3)$, and even worse than the non-BLAS scaling.

Considering now the scaling as a function of the number of threads P , we again find a similar result: that for a fixed number of rows N , the scaling goes as $\mathcal{O}(P^{-1})$, as can be seen in Figure 9. This is once again good news and agrees with what we would expect: the scaling should still be the same as in the non-BLAS case.

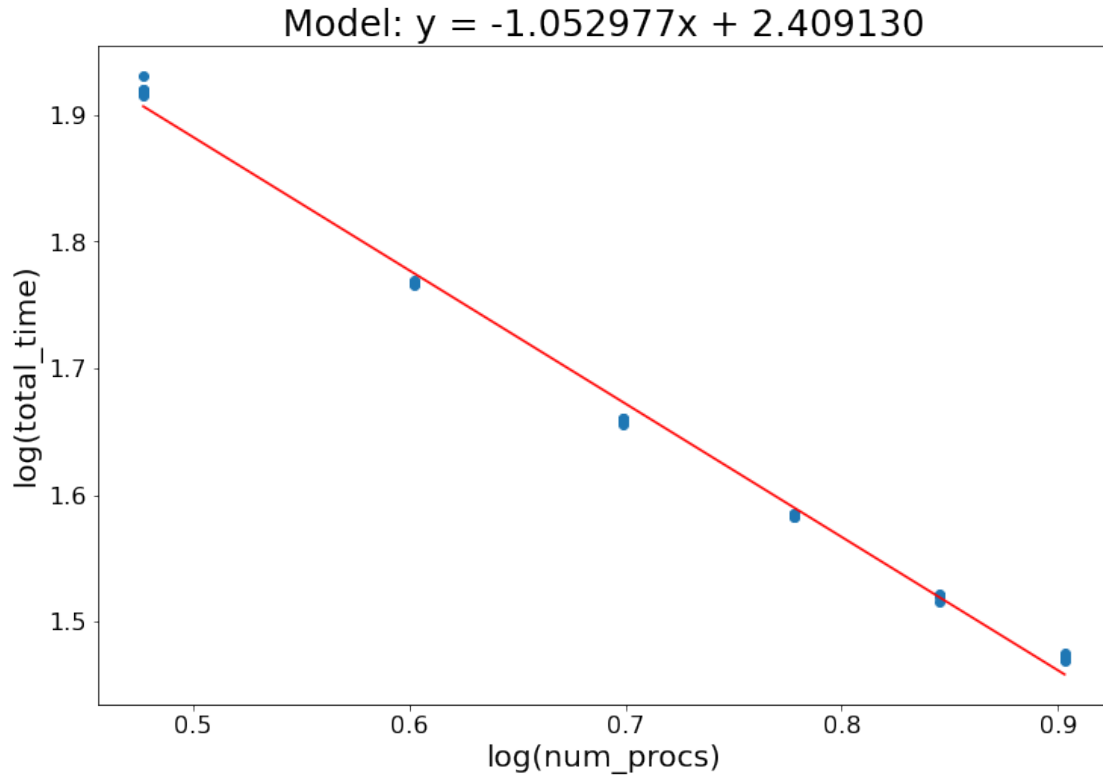


Figure 9: Plotting the runtime of the program as a function of the number of threads P , we find the scaling for a fixed number of rows N (here $N = 3000$) to be of the order $\mathcal{O}(P^{-1})$

Looking at the speedup and efficiency as seen in Table 3, we see that our BLAS code runs faster than the non-BLAS code and is therefore more efficient as well. It also appears that that efficiency of the program increases as we allocate more threads. As a test, I ran the code with $N = 3000$ rows and $P = 12$ threads, and found the efficiency to drop off slightly to 0.59.

Number of Threads	Speedup	Scaling Factor a	Efficiency
3	2.05	0.63	0.63
4	2.65	0.64	0.64
5	3.22	0.65	0.65
6	3.80	0.65	0.65
7	4.35	0.66	0.66
8	4.90	0.66	0.66

Table 3: Summary of the speedup and efficiency of our matrix matrix multiplication code using the BLAS DDOT_ function. These values are taken for a constant number of rows $N = 3000$.

I would expect if we ran the code for say $N = 5000$, the efficiency would be much higher, since if we look at the efficiency for $N = 2000$, as in Table 4, we see that the efficiency of the program nearly

doubles going from $N = 2000$ to $N = 3000$. As a test, I ran the serial code for $N = 5000$, which took 1083.20 seconds. The parallel code ran with $P = 12$ took 128.45 seconds, corresponding to a speedup of 8.43 and an efficiency of 0.70. When run with $P = 8$ instead, it took 188.57 seconds, giving a speedup of 5.74 and an efficiency of 0.72.

Number of Threads	Speedup	Scaling Factor a	Efficiency
3	1.16	0.39	0.39
4	1.52	0.38	0.38
5	1.87	0.37	0.37
6	2.22	0.37	0.37
7	2.56	0.37	0.37
8	2.90	0.36	0.36

Table 4: Summary of the speedup and efficiency of our matrix matrix multiplication code using the BLAS DDOT_ function. These values are taken for a constant number of rows $N = 2000$.

```

1  //////////////////////////////////////
2  // Matrix-Matrix multiplication code AB=C //
3  //////////////////////////////////////
4
5  // Note that I will index arrays from 0 to n-1.
6  // Here workers do all the work and boss just handles collating results
7  // and sending infor about A.
8
9  // include, definitions, globals etc here
10 #include <iostream>
11 #include <fstream>
12 #include <iomanip>
13 #include <random>
14 #include "boost/multi_array.hpp"
15 #include "mpi.h"
16
17 extern "C" {
18     extern double ddot_(int *, double *, int *, double *, int *);
19 }
20
21 using namespace std;
22
23 class MPI_Obj
24 {
25 public:
26     int size;
27     int rank;
28
29     MPI_Obj(int &argc, char **&argv)
30     {
31         MPI_Init(&argc, &argv);
32         MPI_Comm_size(MPI_COMM_WORLD, &size);
33         MPI_Comm_rank(MPI_COMM_WORLD, &rank);
34     }
35
36     ~MPI_Obj()
37     {
38         MPI_Finalize();
39     }
40 };
41
42 void GetArraySize(
43     int &output,
44     int &nrows1,
45     int &nrowscols12,
46     int &ncols2,
47     MPI_Obj &the_mpi,
48     int argc,
49     char **argv
50 ) {

```

```

51     if (the_mpi.rank == 0)
52     {
53         if (argc == 5) {
54             output = atoi(argv[1]);
55             nrow1 = atoi(argv[2]);
56             nrowcols12 = atoi(argv[3]);
57             ncol2 = atoi(argv[4]);
58         } else {
59             cout << "Please enter the number of rows for the first matrix -> ";
60             cin >> nrow1;
61             cout << "Please enter the number of columns/rows for the first/second matrix-> ";
62             cin >> nrowcols12;
63             cout << "Please enter the number of columns for the second matrix-> ";
64             cin >> ncol2;
65         }
66     }
67
68     // send everyone nrow, ncol
69     int buf[3] = {nrow1, nrowcols12, ncol2};
70     MPI_Bcast(buf, 3, MPI_INT, 0, MPI_COMM_WORLD);
71     if (the_mpi.rank != 0)
72     {
73         nrow1 = buf[0];
74         nrowcols12 = buf[1];
75         ncol2 = buf[2];
76     }
77 }
78
79 void SetupArrays(
80     int nrow1,
81     int nrowcols12,
82     int ncol2,
83     boost::multi_array<double, 2> &A,
84     boost::multi_array<double, 2> &B,
85     boost::multi_array<double, 2> &C,
86     vector<double> &Arow,
87     vector<double> &Crow,
88     MPI_Obj &the_mpi
89 ) {
90     uniform_real_distribution<double> unif(-1.0, 10);
91     default_random_engine re;
92
93     B.resize(boost::extents[nrowcols12][ncol2]);
94     Crow.reserve(ncol2); Crow.resize(ncol2); // Main process will need to store the values temporarily
95
96     // Boss part
97     if (the_mpi.rank == 0)
98     {
99         // Set size of A
100         A.resize(boost::extents[nrow1][nrowcols12]);
101         C.resize(boost::extents[nrow1][ncol2]);
102
103         // Initialize A
104         for (int i = 0; i < nrow1; ++i) {
105             for (int j = 0; j < nrowcols12; ++j) {
106                 // Identity
107                 // if (i == j)
108                 //     A[i][j] = 1.0;
109                 // else
110                 //     A[i][j] = 0.0;
111                 // Reverse Identity
112                 // if (i == (nrowcols12 - j - 1))
113                 //     A[i][j] = 1.0;
114                 // else
115                 //     A[i][j] = 0.0;
116                 // Random
117                 A[i][j] = unif(re);
118             }
119         }
120
121         // Initialize B
122         for (int i = 0; i < nrowcols12; ++i) {
123             for (int j = 0; j < ncol2; ++j) {
124                 // Identity
125                 // if (i == j)
126                 //     B[i][j] = 1.0;
127                 // else

```

```

128         // B[i][j] = 0.0;
129         // Reverse Identity
130         // if (i == (ncols2 - j - 1))
131         // B[i][j] = 1.0;
132         // else
133         // B[i][j] = 0.0;
134         // Random
135         // B[i][j] = unif(re);
136         // Other
137         if (i == j)
138             B[i][j] = (double) i + 1;
139         else
140             B[i][j] = 0;
141     }
142 }
143 } else {
144     // Worker part
145     // Allocate space for 1 row of A and 1 row of the answer C
146     Arow.reserve(nrowscols12); Arow.resize(nrowscols12);
147 }
148 MPI_Bcast(&B[0][0], nrowscols12*ncols2, MPI_DOUBLE, 0, MPI_COMM_WORLD);
149 }
150
151 void Output(int output, boost::multi_array<double, 2> &array, MPI_Obj &the_mpi)
152 {
153     if (the_mpi.rank == 0 && output) {
154         cout << endl << fixed << setprecision(4);
155
156         for (int i = 0; i < array.shape()[0]; i++) {
157             for (int j = 0; j < array.shape()[1]; j++) {
158                 cout << array[i][j];
159                 if (j < array.shape()[1] - 1) cout << ", ";
160             }
161             cout << endl;
162         }
163         cout << endl;
164
165         cout << scientific;
166     }
167 }
168
169 int main(int argc, char **argv)
170 {
171     // Data File
172     // ofstream datafile("/home/jared/Desktop/mm-timings.txt", ios_base::app);
173     ofstream datafile("mmblas-timings.txt", ios_base::app);
174
175     // initialize MPI
176     MPI_Obj the_mpi(argc, argv);
177     if (the_mpi.size < 2) MPI_Abort(MPI_COMM_WORLD, 1);
178
179     // determine/distribute size of arrays here
180     int output = 1, nrows1 = 0, nrowscols12 = 0, ncols2 = 0;
181     GetArraySize(output, nrows1, nrowscols12, ncols2, the_mpi, argc, argv);
182     if (the_mpi.size - 1 > nrows1) {
183         MPI_Abort(MPI_COMM_WORLD, -1);
184     }
185
186     boost::multi_array<double, 2> A;
187     boost::multi_array<double, 2> B;
188     boost::multi_array<double, 2> C;
189     vector<double> Arow;
190     vector<double> Crow;
191     SetupArrays(nrows1, nrowscols12, ncols2, A, B, C, Arow, Crow, the_mpi);
192
193     MPI_Status status;
194
195     // Timing variables
196     double calc_time = 0, avg_time, total_time;
197
198     // Boss part
199     if (the_mpi.rank == 0)
200     {
201         total_time = MPI_Wtime();
202         // send one row to each worker tagged with row number, assume size<nrows
203         int rowsent = 1;
204         for (int i = 1; i < the_mpi.size; i++)

```

```

205     {
206         MPI_Send(&A[rowsent - 1][0], nrowcols12, MPI_DOUBLE, i, rowsent, MPI_COMM_WORLD);
207         rowsent++;
208     }
209
210     for (int i = 0; i < nrowsl; i++)
211     {
212         MPI_Recv(&Crow[0], ncols2, MPI_DOUBLE, MPI_ANY_SOURCE, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
213         int sender = status.MPI_SOURCE;
214         int row = status.MPI_TAG - 1;
215         memcpy(&C[row][0], &Crow[0], ncols2 * sizeof(double));
216
217         if (rowsent - 1 < nrowsl) {
218             // send new row
219             MPI_Send(&A[rowsent - 1][0], nrowcols12, MPI_DOUBLE, sender, rowsent, MPI_COMM_WORLD);
220             rowsent++;
221         } else {
222             // tell sender no more work to do via a 0 TAG
223             MPI_Send(MPI_BOTTOM, 0, MPI_DOUBLE, sender, 0, MPI_COMM_WORLD);
224         }
225     }
226 }
227 // Worker part: compute dot products of Arow.b until done message recieved
228 else
229 {
230     // Get a row of A
231     MPI_Recv(&Arow[0], nrowcols12, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
232     while (status.MPI_TAG != 0)
233     {
234         for (int i = 0; i < ncols2; i++) {
235             // work out Crow = Arow.B
236             calc_time -= MPI_Wtime();
237             Crow[i] = ddot_(&nrowcols12, &Arow[0], new int(1), &B[0][i], &ncols2);
238             calc_time += MPI_Wtime();
239         }
240         // Send answer of Arow.B back to boss and get another row to work on
241         MPI_Send(&Crow[0], ncols2, MPI_DOUBLE, 0, status.MPI_TAG, MPI_COMM_WORLD);
242         MPI_Recv(&Arow[0], nrowcols12, MPI_DOUBLE, 0, MPI_ANY_TAG, MPI_COMM_WORLD, &status);
243     }
244 }
245 }
246
247 if (the_mpi.rank == 0)
248     total_time = MPI_Wtime() - total_time;
249
250 MPI_Reduce(&calc_time, &avg_time, 1, MPI_DOUBLE, MPI_SUM, 0, MPI_COMM_WORLD);
251 avg_time /= (the_mpi.size - 1); // Boss node doesn't do any of the calculations
252
253 if (the_mpi.rank == 0)
254 {
255     printf("Average calculation time = %.10f\n", avg_time);
256     printf("Total time = %.10f\n", total_time);
257
258     datafile << fixed << setprecision(10);
259     datafile << nrowsl << " " << nrowcols12 << " " << ncols2 << " " <<
260         the_mpi.size << " " << avg_time << " " << total_time << endl;
261 }
262
263 // output c here on Boss node
264 Output(output, C, the_mpi);
265 }

```

4 Solving Poisson's Equation Using LAPACK

The last part of the assignment is a brief introduction to the LAPACK library. The code being tested solves the Poisson equation $\nabla^2 U(\vec{x}) = f(\vec{x})$. In Figure 10, I have run the provided code and plotted the solution using plotly in Python.

The code works by solving a linear system of equations of the form $Ax = b$, where A is a matrix, while x and b are vectors. The program is made more efficient by noting for this problem that the matrix A will be a banded, negative-definite matrix, such that we can use the optimized LAPACK functions `DPBTRF_` and `DPBTRS_` to factor and solve the system.

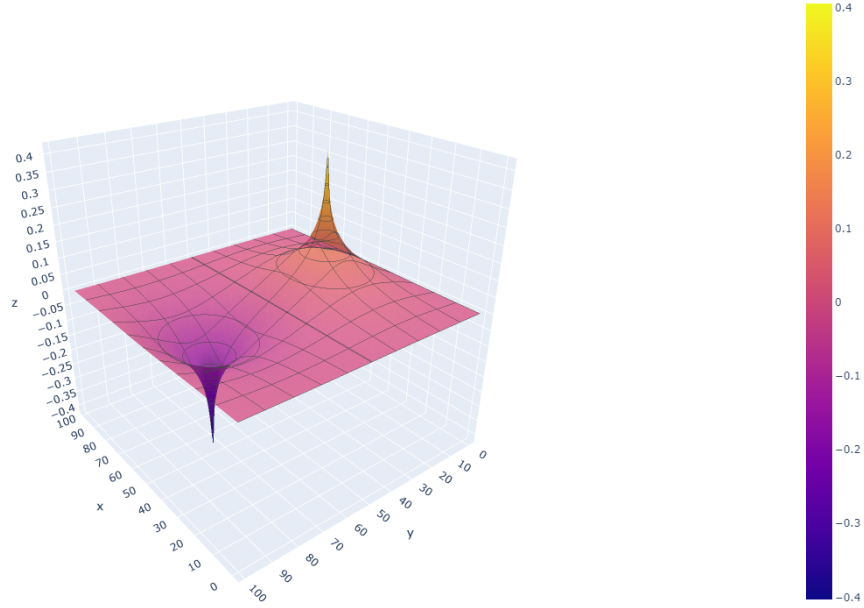


Figure 10: Plot of the solution using $N = 100$ sites.

Testing the run time of the code as a function of the number of sites N , the scaling is found to be of the order $\mathcal{O}(N^{3.8})$, as can be seen in Figure 11. The scaling can be expected to be $\mathcal{O}(N^4)$, as performing a PLU factorization scales as $\mathcal{O}(N^3)$, and back-substitution should contribute a factor proportional to $\mathcal{O}(N)$. Given we are using BLAS routines which are optimized for banded and positive definite matrices, it is convincing our scaling should be slightly less than the worst case scenario of $\mathcal{O}(N^4)$.

Using the linear regression model (red line) in the plot, we can estimate that the largest system that can be solved (on my hardware) in 5 minutes to be roughly $N = 975$. This seems reasonable, since the runtime for a system with $N = 1000$ was roughly 385 seconds.

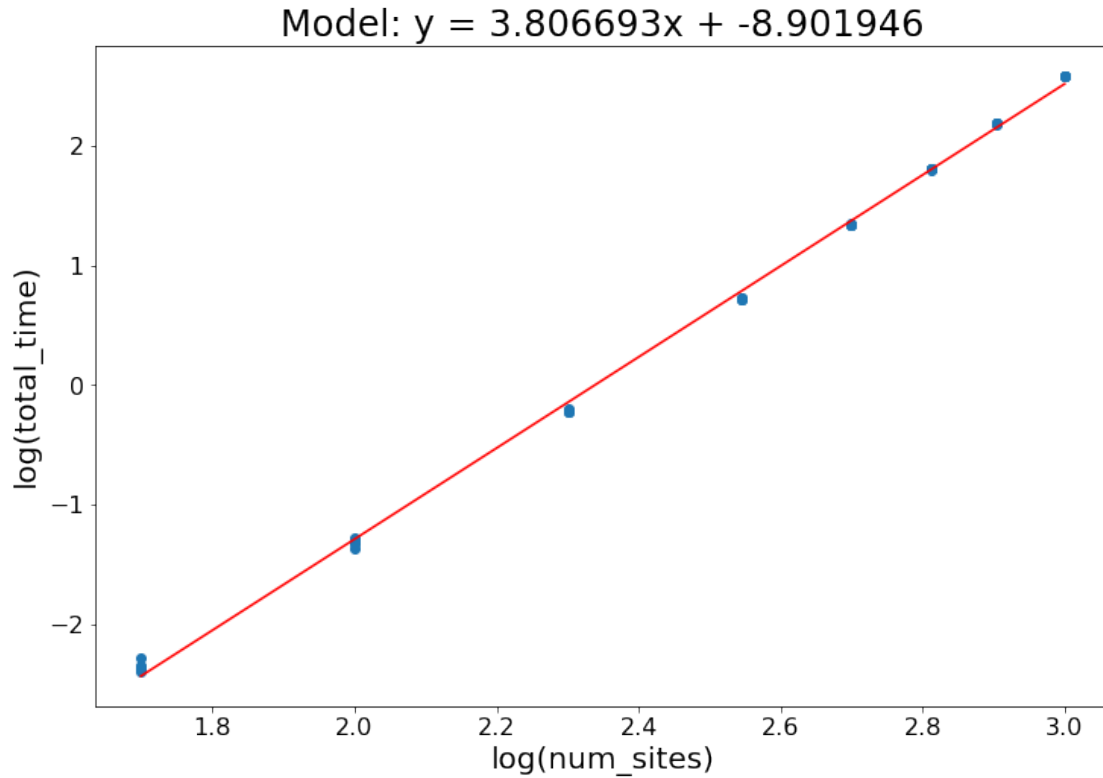


Figure 11: Plotting the runtime of the program as a function of the number of sites N (on a log-log scale), we find that the total runtime scales as $\mathcal{O}(N^{3.8})$.

```

1 // Program to compute solution to -laplacian u = -F
2 // compile with
3 // run with mpirun -n 1 <executable> output_cout output_file N
4
5 #include <iostream>
6 #include <iomanip>
7 #include <fstream>
8 #include <mpi.h>
9 #include "boost/multi_array.hpp"
10
11 // LAPACK library files
12 extern "C"
13 {
14     // general band factorize routine
15     extern int dpbtrf_(char *, int *, int *, double *, int *, int *);
16     // general band solve using factorization routine
17     extern int dpbtrs_(char *, int *, int *, int *, double *, int *, double *, int *, int *);
18 }
19
20 void AbInit(int N, boost::multi_array<double, 2> &Ab)
21 {
22     // Coefficient Matrix: initialize
23     for (int i = 0; i < N * N; i++)
24         for (int j = 0; j < N + 1; j++)
25         {
26             Ab[j][i] = 0.0;
27             if (j == 0)
28                 Ab[j][i] = -1.0;
29             if (j == N - 1 && i % N)
30                 Ab[j][i] = -1.0;
31             if (j == N)
32                 Ab[j][i] = 4.0;
33         }
34     // std::cout << "Ab initialized \n";
35     // Printout Ab for testing, small N only

```

```

36     // for (int i=0; i < N+1; i++) {
37     //     std::cout << Ab[i][0];
38     //     for (int j=1; j < N*N; j++) {
39     //         std::cout << " " << Ab[i][j];
40     //     }
41     //     std::cout << "\n";
42     // }
43 }
44
45 void RHSInitialize(int N, std::vector<double> &F, double bcx, double bcy)
46 {
47     // RHS: fill in boundary condition values
48     for (int i = 0; i < N; i++)
49     {
50         F[i] += bcx; // bottom boundary
51         F[N * N - i - 1] += bcx; // top boundary
52         F[i * N] += bcy; // left boundary
53         F[i * N + N - 1] += bcy; // right boundary
54     }
55     // std::cout << "RHS initialized\n";
56
57     // RHS: fill in actual right-hand side, some "charges", actually h^2*charge
58     F[N / 4 * N + N / 2] += 0.5;
59     F[3 * N / 4 * N + N / 2] += -0.5;
60 }
61
62 int main(int argc, char** argv)
63 {
64     double time = MPI_Wtime();
65     // ofstream timingfile("/home/jared/Desktop/lp-timings.txt", ios_base::app);
66     std::ofstream timingfile("lp-timings.txt", std::ios_base::app);
67
68     // ofstream datafile("/home/jared/Desktop/lp-data.txt");
69     std::ofstream datafile("lp-data.txt");
70
71     // Coefficient Matrix: declare
72     int output_cout = 1, output_file = 0;
73     int N = 10;
74     if (argc == 4) {
75         output_cout = atoi(argv[1]);
76         output_file = atoi(argv[2]);
77         N = atoi(argv[3]);
78     }
79
80     int M = N + 1;
81     int ABcols = N * N;
82     boost::multi_array<double, 2> Ab(boost::extents[M][ABcols], boost::fortran_storage_order());
83
84     AbInit(N, Ab); // Initialize coefficient matrix
85
86     // Coefficient Matrix: factorize
87     char uplo = 'U';
88     int KD = N;
89     int info;
90     dpbtrf_(&uplo, &ABcols, &KD, &Ab[0][0], &M, &info);
91     if (info)
92     {
93         std::cout << "Ab failed to factorize, info = " << info << "\n";
94         exit(1);
95     }
96
97     // RHS: declare
98     const double bcx = 0.0, bcy = 0.0; // boundary conditions along x and y assume same on both sides
99     std::vector<double> F(N * N, 0.0);
100
101     RHSInitialize(N, F, bcx, bcy); // set up boundary conditions and right hand side
102
103     // Solve system
104     int Bcols = 1;
105     dpbtrs_(&uplo, &ABcols, &KD, &Bcols, &Ab[0][0], &M, &F[0], &ABcols, &info);
106     if (info)
107     {
108         std::cout << "System solve failed, info = " << info << "\n";
109         exit(1);
110     }
111
112     time = MPI_Wtime() - time;

```

```

113     timingfile << N << " " << time << std::endl;
114
115     // Output solution
116     if (output_cout) {
117         for (int i = 0; i < N; i++)
118             {
119                 std::cout << F[i * N];
120                 for (int j = 1; j < N; j++)
121                     {
122                         std::cout << " " << F[i * N + j];
123                     }
124                 std::cout << "\n";
125             }
126     }
127     if (output_file) {
128         datafile << std::fixed << std::setprecision(10);
129         for (int i = 0; i < N; i++)
130             {
131                 datafile << F[i * N];
132                 for (int j = 1; j < N; j++)
133                     {
134                         datafile << " " << F[i * N + j];
135                     }
136                 datafile << "\n";
137             }
138     }
139
140     return 0;
141 }

```


A Jupyter Notebook

Below is an example notebook I have used in order to generate and analyze my timing data. The notebook will compile and run the given program a number of times for various numbers of threads and matrix dimensions. It then creates linear regression models from the data which can be used to determine the speedup and efficiency of parallelization.

plots

November 21, 2022

```
[ ]: import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import pandas as pd
import subprocess as sp
import itertools as it

[ ]: def genData(procs, sizes, data_file, cycles=20, clean=False):
    sp.run(
        "/usr/bin/mpiCC -O2 -o 15-matrix-vector.o ./10-matrix-vector.cpp",
        shell=True,
        stdout=sp.DEVNULL,
        stderr=sp.DEVNULL
    )
    if clean:
        sp.run(
            f'echo "num_rows num_cols num_procs avg_time total_time" >{
↪{data_file}',
            shell=True,
            stdout=sp.DEVNULL,
            stderr=sp.DEVNULL
        )

    for size, proc in it.product(sizes, procs):
        for _ in range(cycles):
            sp.run(
                f"/usr/bin/mpirun -n {proc} ./15-matrix-vector.o 0 {size}{
↪{size}",
                shell=True,
                stdout=sp.DEVNULL,
                stderr=sp.DEVNULL
            )
            print(f"Finished {size}x{size} with {proc} processes")

def plotModel(
    data: pd.DataFrame,
    xdata: str,
```

```

ydata: str,
plot: bool,
loglog: bool
):
    x = data[xdata].values.reshape(-1, 1)
    y = data[ydata].values.reshape(-1, 1)
    xname = xdata
    yname = ydata
    if loglog:
        x = np.log10(x)
        y = np.log10(y)
        xname = "log(" + xdata + ")"
        yname = "log(" + ydata + ")"

    linear_model = LinearRegression()
    linear_model.fit(x, y)
    model = linear_model.predict(x)

    if plot:
        fig, ax = plt.subplots(figsize=(12,8))
        ax.set_title("Model: y = {:.6f}x + {:.6f}".format(linear_model.
↪coef_[0][0], linear_model.intercept_[0]), size=24)
        ax.set_xlabel(xname, size=20)
        ax.tick_params(axis='x', labelsize=16)
        ax.set_ylabel(yname, size=20)
        ax.tick_params(axis='y', labelsize=16)
        ax.scatter(x, y)
        ax.plot(x, model, color='red')
        plt.show()

    return linear_model

```

```

[ ]: # Variables
# data_file = "/home/jared/Desktop/mv-timings.txt"
data_file = "./mv-timings.txt"
procs = np.array([3, 4, 5, 6, 7, 8])
sizes = np.array([10, 50, 100, 250, 500, 1000, 10000])

# Generate data if needed
# genData(procs=procs, sizes=sizes, data_file=data_file, clean=True)

# Load the data
data = pd.read_csv(data_file, delimiter=" ")

```

```

[ ]: num_procs = 4
num_rows = 10000
min_procs = 2

```

```

min_rows = 100

data_reduced_row_scaling = data[data["num_procs"] == num_procs]
data_reduced_row_scaling = data_reduced_row_scaling[data_reduced_row_scaling["num_rows"] > min_rows]

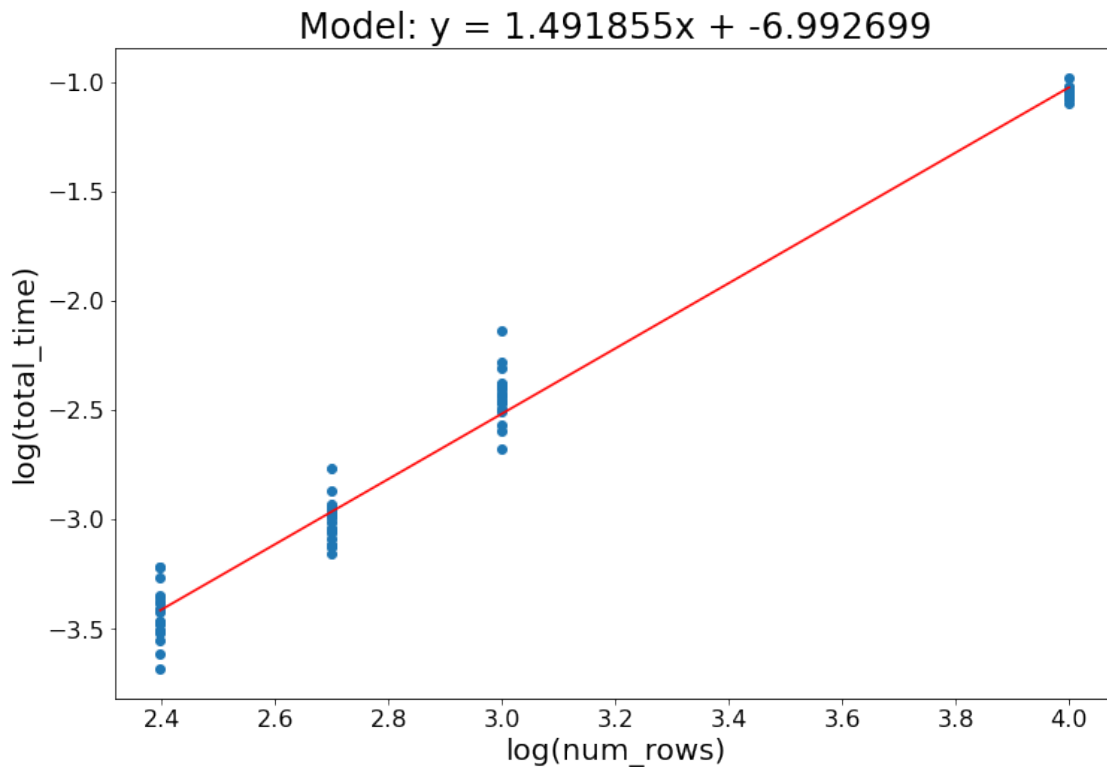
data_reduced_proc_scaling = data[data["num_rows"] == num_rows]
data_reduced_proc_scaling = data_reduced_proc_scaling[data_reduced_proc_scaling["num_procs"] > min_procs]

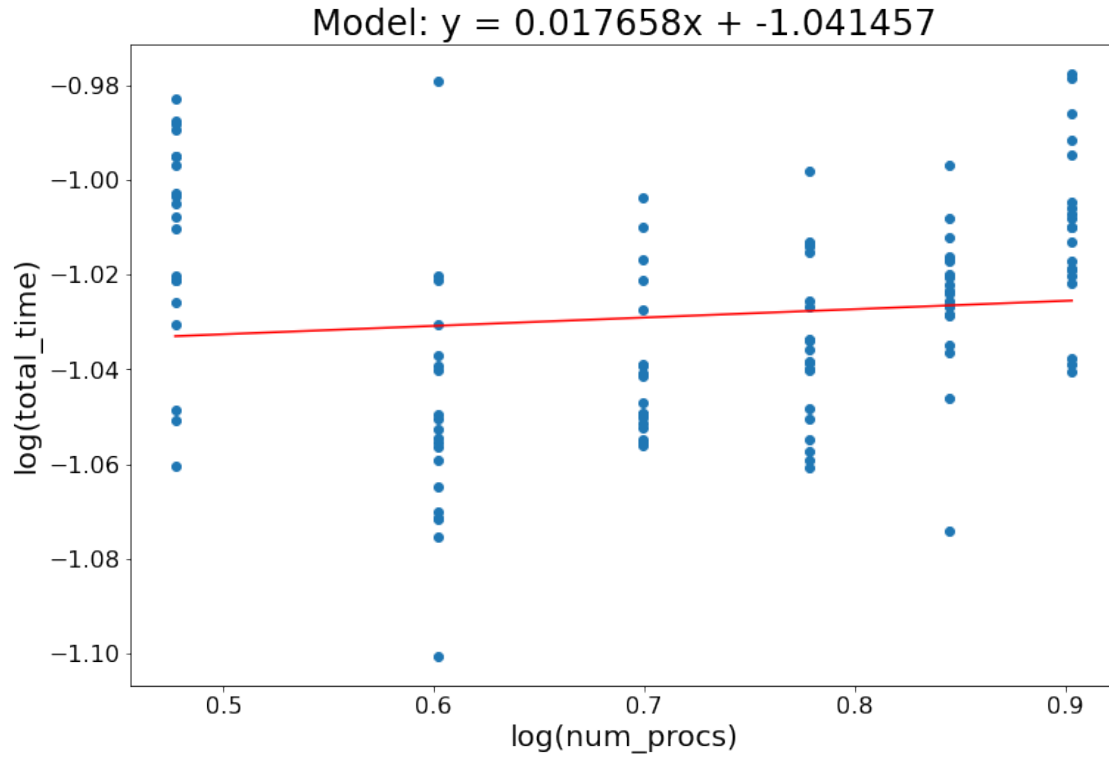
```

```

[ ]: reduced_row_scaling_model = plotModel(data_reduced_row_scaling, "num_rows", "total_time", plot=True, loglog=True)
reduced_proc_scaling_model = plotModel(data_reduced_proc_scaling, "num_procs", "total_time", plot=True, loglog=True)

```





```
[ ]: num_rows = 10000

serial_times = {
    10: 0.0000047090,
    50: 0.0000092870,
    100: 0.0000360770,
    250: 0.0001662810,
    500: 0.0008027130,
    1000: 0.0025987250,
    10000: 0.4020480850,
}

prediction = np.array([
    10**val[0] for val in
    reduced_proc_scaling_model.predict(procs.reshape(-1, 1))
])

speedup = serial_times[num_rows] / prediction
print("speedup = \n", np.array_str(speedup))

a = speedup / procs
print("a = \n", np.array_str(a))
```

```
efficiency = speedup / procs
print("efficiency = \n", np.array_str(efficiency))
```

```
speedup =
[3.91524996 3.75924844 3.60946274 3.46564519 3.32755799 3.19497282]
a =
[1.30508332 0.93981211 0.72189255 0.57760753 0.47536543 0.3993716 ]
efficiency =
[1.30508332 0.93981211 0.72189255 0.57760753 0.47536543 0.3993716 ]
```