

plots

November 21, 2022

```
[ ]: import numpy as np
from sklearn.linear_model import LinearRegression
import matplotlib.pyplot as plt
import pandas as pd
import subprocess as sp
import itertools as it

[ ]: def genData(procs, sizes, data_file, cycles=20, clean=False):
    sp.run(
        "/usr/bin/mpiCC -O2 -o 15-matrix-vector.o ./10-matrix-vector.cpp",
        shell=True,
        stdout=sp.DEVNULL,
        stderr=sp.DEVNULL
    )
    if clean:
        sp.run(
            f'echo "num_rows num_cols num_procs avg_time total_time" >{data_file}',
            shell=True,
            stdout=sp.DEVNULL,
            stderr=sp.DEVNULL
        )

    for size, proc in it.product(sizes, procs):
        for _ in range(cycles):
            sp.run(
                f"/usr/bin/mpirun -n {proc} ./15-matrix-vector.o 0 {size} {size}",
                shell=True,
                stdout=sp.DEVNULL,
                stderr=sp.DEVNULL
            )
            print(f"Finished {size}x{size} with {proc} processes")

def plotModel(
    data: pd.DataFrame,
    xdata: str,
```

```

        ydata: str,
        plot: bool,
        loglog: bool
    ):
        x = data[xdata].values.reshape(-1, 1)
        y = data[ydata].values.reshape(-1, 1)
        xname = xdata
        yname = ydata
        if loglog:
            x = np.log10(x)
            y = np.log10(y)
            xname = "log(" + xdata + ")"
            yname = "log(" + ydata + ")"

        linear_model = LinearRegression()
        linear_model.fit(x, y)
        model = linear_model.predict(x)

        if plot:
            fig, ax = plt.subplots(figsize=(12,8))
            ax.set_title("Model: y = {:.6f}x + {:.6f}".format(linear_model.
↪coef_[0][0], linear_model.intercept_[0]), size=24)
            ax.set_xlabel(xname, size=20)
            ax.tick_params(axis='x', labelsize=16)
            ax.set_ylabel(yname, size=20)
            ax.tick_params(axis='y', labelsize=16)
            ax.scatter(x, y)
            ax.plot(x, model, color='red')
            plt.show()

        return linear_model

```

```

[ ]: # Variables
# data_file = "/home/jared/Desktop/mv-timings.txt"
data_file = "./mv-timings.txt"
procs = np.array([3, 4, 5, 6, 7, 8])
sizes = np.array([10, 50, 100, 250, 500, 1000, 10000])

# Generate data if needed
# genData(procs=procs, sizes=sizes, data_file=data_file, clean=True)

# Load the data
data = pd.read_csv(data_file, delimiter=" ")

```

```

[ ]: num_procs = 4
      num_rows = 10000
      min_procs = 2

```

```

min_rows = 100

data_reduced_row_scaling = data[data["num_procs"] == num_procs]
data_reduced_row_scaling = data_reduced_row_scaling[data_reduced_row_scaling["num_rows"] > min_rows]


data_reduced_proc_scaling = data[data["num_rows"] == num_rows]
data_reduced_proc_scaling = data_reduced_proc_scaling[data_reduced_proc_scaling["num_procs"] > min_procs]

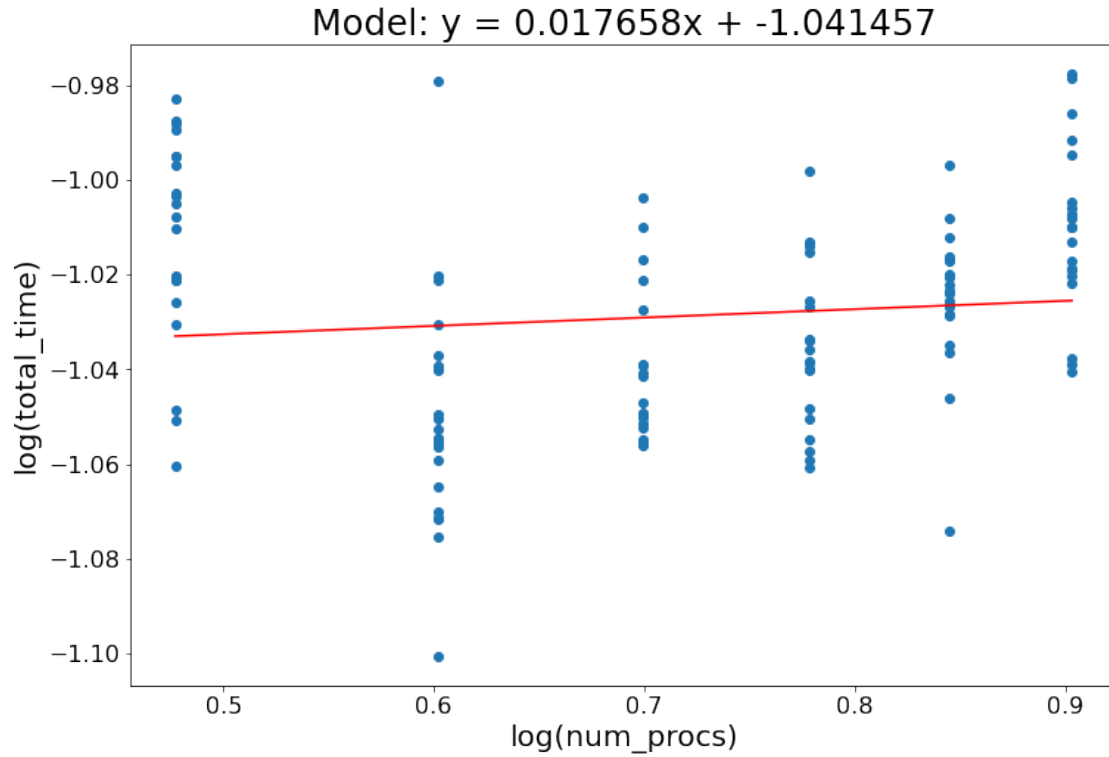
```

```

[ ]: reduced_row_scaling_model = plotModel(data_reduced_row_scaling, "num_rows", "total_time", plot=True, loglog=True)
reduced_proc_scaling_model = plotModel(data_reduced_proc_scaling, "num_procs", "total_time", plot=True, loglog=True)

```





```
[ ]: num_rows = 10000

serial_times = {
    10: 0.0000047090,
    50: 0.0000092870,
    100: 0.0000360770,
    250: 0.0001662810,
    500: 0.0008027130,
    1000: 0.0025987250,
    10000: 0.4020480850,
}

prediction = np.array([
    10**val[0] for val in
    reduced_proc_scaling_model.predict(procs.reshape(-1, 1))
])

speedup = serial_times[num_rows] / prediction
print("speedup = \n", np.array_str(speedup))

a = speedup / procs
print("a = \n", np.array_str(a))
```

```
efficiency = speedup / procs
print("efficiency = \n", np.array_str(efficiency))
```

```
speedup =
[3.91524996 3.75924844 3.60946274 3.46564519 3.32755799 3.19497282]
a =
[1.30508332 0.93981211 0.72189255 0.57760753 0.47536543 0.3993716 ]
efficiency =
[1.30508332 0.93981211 0.72189255 0.57760753 0.47536543 0.3993716 ]
```