

Intro to Parallel Computing


January 14, 2014 7:02 PM

Announcements

- course outline is on OWL, make sure you have access
- we will use C/C++ and MPI in this course so you need access to a computer running these
- a Linux-like environment is needed
- WSL (Windows Subsystem for Linux) should work for program development if you install the MPI runtime & developer packages in setup.
- on Macs, you can probably get by using the XCode package. There is also a MPI implementation you will need
- there are some tips on OWL in the "Set-Up" file
- you need to get a ShareNet account
→ see ShareNet website
(a line terminal is needed to access ShareNet computers)

Lecture 1

Back in the day
computers looked
like →
(not really...)



So you
could play
"Mario Chariot"

1980	1995	2005	2015
comodore vic 20	gateway pentium PC	Dell pentium 4 PC	lenovo i7 PC
CPU 1 MHz	100 MHz	2 GHz	4x 3 GHz
Mem 5 kByte RAM 20 kByte ROM	4 Mbyte	1 GByte	4 GByte
			GPU 1536x 1GHz (nvidia geforce 770)

$\times 100$

$\times 200$

$\times 1.5$

this is the bigger CPU gain
i.e. multicore processor

This is by far
the biggest
improvement

Reasons to parallelize
→ time (clock time)
→ memory

Classification of parallelization according to data & instruction schemes

	same data	multiple data
single instruction	SISD (serial)	SIMD
multiple instructions	MISD	MIMD

note: different computer architecture may be required for different schemes but some architectures can run different schemes

eg ① SISD

Compute $F_i = \sum_n f_n$

Algorithm

$$\begin{array}{l}
 F_i = f_1 \\
 F_i += f_2 \\
 F_i += f_3 \\
 \vdots \\
 F_i += f_n
 \end{array}
 \quad \begin{array}{l} \text{time} \\ \downarrow \end{array}$$

eg ② SIMD

Compute $A = 4 \cdot B$ where A, B $1 \times n$ arrays
n divisible by 3

Algorithm


	CPU1	CPU2	CPU3
time ↓	$A[1] = 4 * B[1]$	$A[2] = 4 * B[2]$	$A[3] = 4 * B[3]$
	$A[4] = 4 * B[4]$	$A[5] = 4 * B[5]$	$A[6] = 4 * B[6]$
	\vdots	\vdots	\vdots
	$A[n-2] = 4 * B[n-2]$	$A[n-1] = 4 * B[n-1]$	$A[n] = 4 * B[n]$

The division of data is clearly not unique. Different divisions may result in faster/slower code.

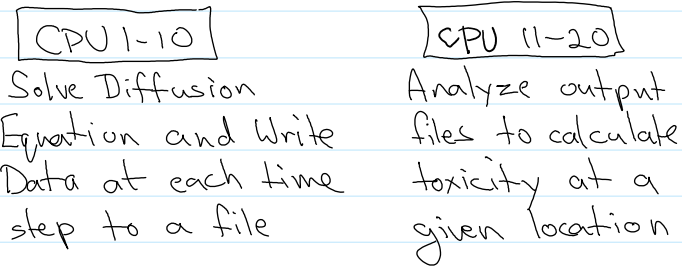
eg ③ MISD

Compute $F_n = \frac{1}{N} \sum_n f_n^N$ (moments problem)

	CPU1	CPU2	...	CPU _N
time ↓	$F_1 = f_1$	$F_2 = f_1^2$		$F_N = f_1^N$
	$F_1 += f_2$	$F_2 += f_2^2$		$F_N += f_2^N$
	$F_1 += f_3$	\vdots		\vdots
	\vdots			




eg ④ MIMD different CPUs have different tasks & act on different data sets




Classification of parallel computing according to memory

Shared Memory (typical desktop today)



Distributed Memory (beowulf cluster)



In practise, most systems today are some combination. However, it is useful to develop algorithms based on a single paradigm in order to make portable code.

MPI is based on a distributed memory paradigm. It provides a library of functions to facilitate communication between processors

$$\text{MPI} = \begin{cases} \text{Message} \\ \text{Passing} \\ \text{Interface} \end{cases}$$

MPI

- library containing ≈ 100 functions
- ~ 6 are really important:
 - 1) initialize MPI
 - 2) find # of processes (specified when program submitted)
 - 3) find out which process I am

main
point
⇒

4) send a message

⇒ 5) receive a message

6) terminate MPI (essential. Make sure

you always do this! MPI is C based not C++)

Eg. "Hello World" program in C (source code is on OWL)

```
#include <stdio.h>
#include "mpi.h"
main(int argc, char** argv)
{
    int nprocs, nid;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
    MPI_Comm_rank(MPI_COMM_WORLD, &nid);

    printf("Hello from processor %i of %i\n", nid, nprocs);

    MPI_Finalize();
}
```

MPI header files (always need this)

starts MPI and code communications on all processors

runs on each node independently

close of all the processes

comments:

MPI_Init must be called (exactly) once before all other MPI commands


MPI_Finalize should be called after all other MPI commands

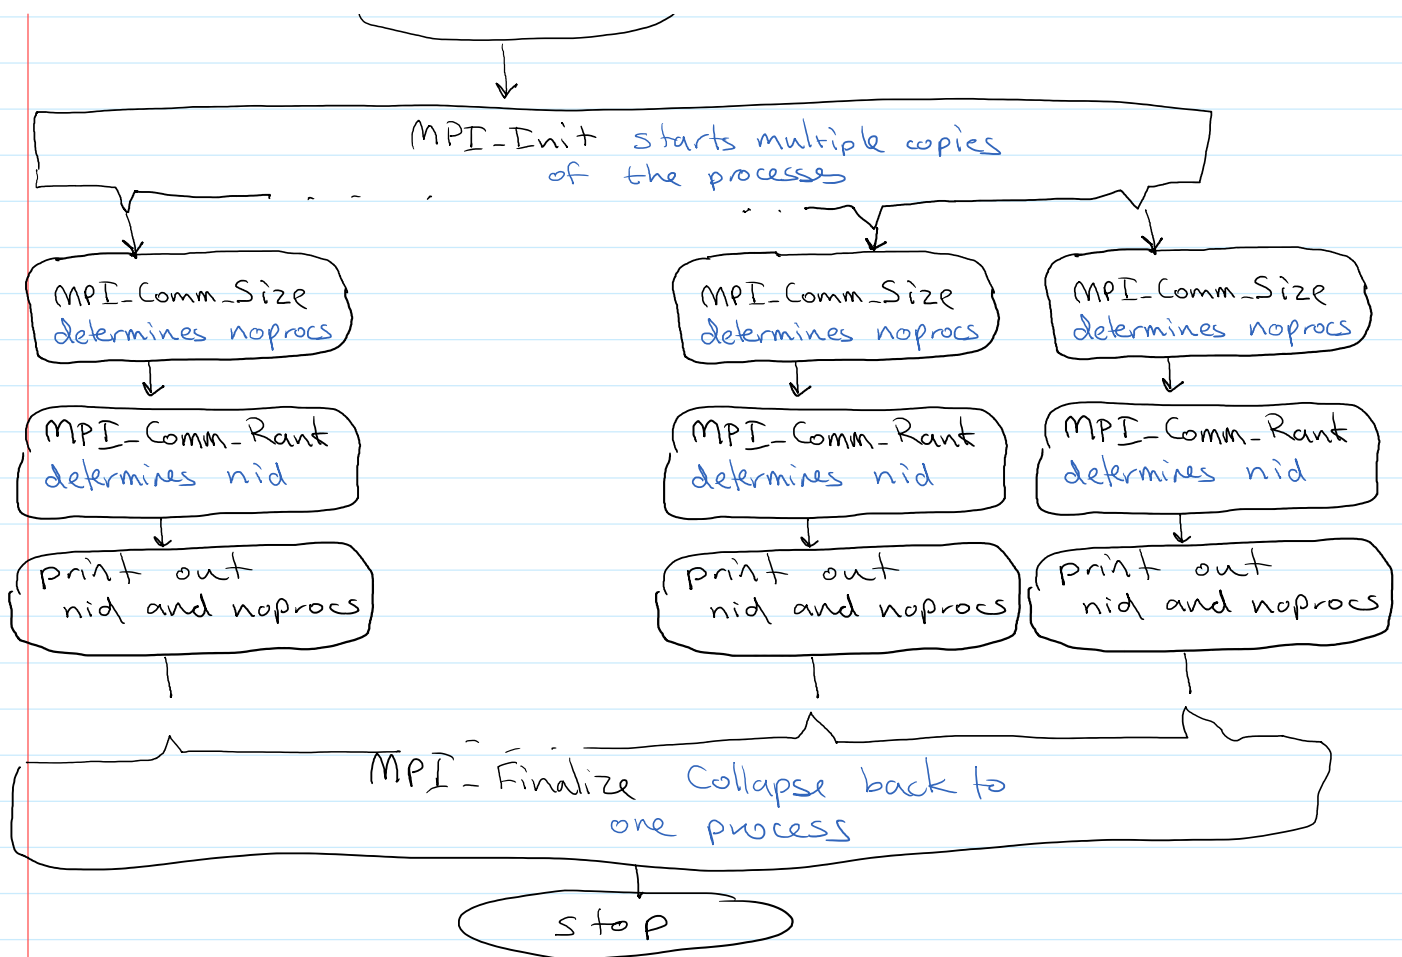
MPI_Comm_size finds out number of processors (put into *nprocs*)

MPI_Comm_rank finds out which process I am
 $nid \in \{0, \dots, nprocs - 1\}$

MPI_COMM_WORLD is a MPI "communicator"

A flow-chart of what happens:





Note that once multiprocesses are running they are independent and variables like "nid" are different for each process.

printf statement is executed on each processor independently (as are previous 2 commands) and there is a different local version of nid and nprocs on each processor, potentially with different values

The different processes are not typically in sync so the output may come in any order and be different if you run the code again.

Typical compile line is something like

```
mpicc -o hello.o hello.c
```

Typical run command is

```
mpirun -np 4 specify # of processes ./hello.o
```

Expected Output:

Hello from processor 0 of 4	} perhaps not in this order
Hello from processor 1 of 4	
Hello from processor 2 of 4	
Hello from processor 3 of 4	

MPI is a library written in C. It also has a Fortran version interface. There was a long-term effort to get a C++ version but it was buggy, largely shunned in favor of the C version and was abandoned as unnecessary (C++ can just use the C library).

For C++ it is always a good idea to try to follow the RAII paradigm and try to automate cleanup using class destructors to take care of it. To make our C and C++ code as similar as possible we will first rewrite hello.c to put MPI variables into a data structure and add a startup function:

```
#include <stdio.h>
#include "mpi.h"
```

```
struct mpi_vars { ← data structure type declaration
    int NProcs;
    int MyID;
};
```

```
struct mpi_vars mpi_start(int argc, char** argv)
{
    struct mpi_vars this_mpi;

    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &this_mpi.NProcs);
    MPI_Comm_rank(MPI_COMM_WORLD, &this_mpi.MyID);

    return this_mpi;
}
```

```
int main(int argc, char** argv)
{
    struct mpi_vars this_mpi = mpi_start(argc, argv);
    printf("Hello from processor %i of %i\n",
           this_mpi.MyID, this_mpi.NProcs);
}
```

← initialize MPI data structure using start function

```

MPI_Finalize();

return 0;
}

```

- Note:
1. Only one instance of the mpi-vars should be created (on each MPI thread)
 2. The mpi-start function should only ever be called once.
 3. In C, MPI_Finalize() should always be called before you exit the program

Translation of this to C++ is then very straightforward:

```

#include <iostream>
#include "mpi.h"

```

```

class MPI_stuff
{

```

```

public:

```

```

    int NProcs;
    int MyID;

```

} our mpi info

```

MPI_stuff(int &argc, char** &argv) ← put MPI initialization
{                                     into class constructor
    MPI_Init(&argc, &argv);
    MPI_Comm_size(MPI_COMM_WORLD, &NProcs);
    MPI_Comm_rank(MPI_COMM_WORLD, &MyID);
}

```

```

~MPI_stuff() ← put MPI cleanup
{             into class destructor
    MPI_Finalize();
}
};

```

```

int main(int argc, char** argv)
{

```

```

    MPI_stuff the_mpi(argc, argv); ← MPI will be started automatically
                                   by the MPI_stuff constructor

```

```

    std::cout << "Hello from processor " << the_mpi.MyID << " of "
               << the_mpi.NProcs << std::endl;

```

```

    return 0; ← the_mpi variable goes out of scope here
               all the called to clean off MPI

```

```
return 0; ← the mpi variable goes out of scope here  
}          so the destructor will be called to close off MPI
```

- Note:
- 1) The class MPI_stuff should only ever have one instance in existence (but this one instance may/will have different values for MyID on each processor)
 - 2) MPI_Finalize will automatically be called when the mpi variable goes out of scope (useful if program exits in unexpected manner too).
 - 3) Compilation is typically something like
mpicc -o hello.o hello.cpp
note
capital c's here

"Homework"

1. Sign up for a ShareNet account and online "training seminar"
2. Try to get Hello world running on your machine