

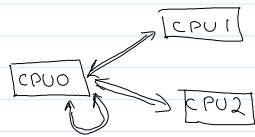
Message Passing between Processes

January 14, 2014 9:43 PM

Passing Data

We will look at 2 different "modes":

- i) One-to-all (broadcast model)
- ii) All-to-one (reverse of broadcast, a "reduce")



We will use the broadcast model in our 1st example.

$$\text{eg. } \ln 2 = \sum_{n=1}^{\infty} (-1)^{n+1} \frac{1}{n} = \sum_{n=0}^{N-1} (-1)^n \frac{1}{n+1}$$

assume "big"

shift sum to start with n=0

divide up sum amongst processors

$$= \sum_{nid=0}^{Nprocs-1} \left(\sum_{i=nid Nprocs}^N (-1)^i \frac{1}{i+1} \right)$$

increment for sum

So cpu 0 computes $(-1)^0 \frac{1}{0+1} + (-1)^{Nprocs} \frac{1}{Nprocs+1} + (-1)^{2Nprocs} \frac{1}{2Nprocs+1} \dots$

cpu 1 computes $(-1)^1 \frac{1}{1+1} + (-1)^{Nprocs+1} \frac{1}{(Nprocs)+1} + (-1)^{2Nprocs+1} \frac{1}{(2Nprocs)+1} + \dots$

:

cpu ($Nprocs-1$) computes $(-1)^{Nprocs-1} \frac{1}{((Nprocs)-1)+1} + (-1)^{2Nprocs-1} \frac{1}{(2(Nprocs)-1)+1}$

- Let's use a Boss (typically node or cpu 0) - worker (all the rest) model
- Boss does I/O, everyone works on a different set of terms in the sum $n=nid, nid+Nprocs, \text{etc}$
- Results from each processor need to be collected and added together to get final result.

Here is some code that does this:

```

#include <iostream>
#include "mpi.h"

class MPI_stuff
{
public:
    int NProcs;
    int MyID;

    MPI_stuff(int &argc, char** &argv)
    {
        MPI_Init(&argc, &argv);
        MPI_Comm_size(MPI_COMM_WORLD, &NProcs);
        MPI_Comm_rank(MPI_COMM_WORLD, &MyID);
    }

    ~MPI_stuff()
    {
        MPI_Finalize();
    }
};

float compute_lnsum(const MPI_stuff &the_mpi, const int N)
{
    float sum = 0, Gsum;

    for(int i = the_mpi.MyID; i < N; i += the_mpi.NProcs)
        if(the_mpi.MyID % 2)
            sum -= (float) 1 / (i + 1);
        else
            sum += (float) 1 / (i + 1);

    MPI_Reduce(&sum, &Gsum, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD);
    ← collect & sum result from each processor
    return Gsum;
}

int main(int argc, char** argv)
{
    MPI_stuff the_mpi(argc, argv); ← setup MPI

    int N;
    if(the_mpi.MyID == 0){ ← only one process (Boss) should interact
        with user
        std::cout << "Please enter the number of terms N -> ";
        std::cin >> N;
    }
    MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD); ← send/recieve N via a
    broadcast

    float Gsum = compute_lnsum(the_mpi, N); ← compute sum in parallel
    (each worker & Boss calls this)

    if(the_mpi.MyID == 0) ← only Boss outputs result (to avoid duplication)
        std::cout << "An estimate of ln(2) is " << Gsum << std::endl;

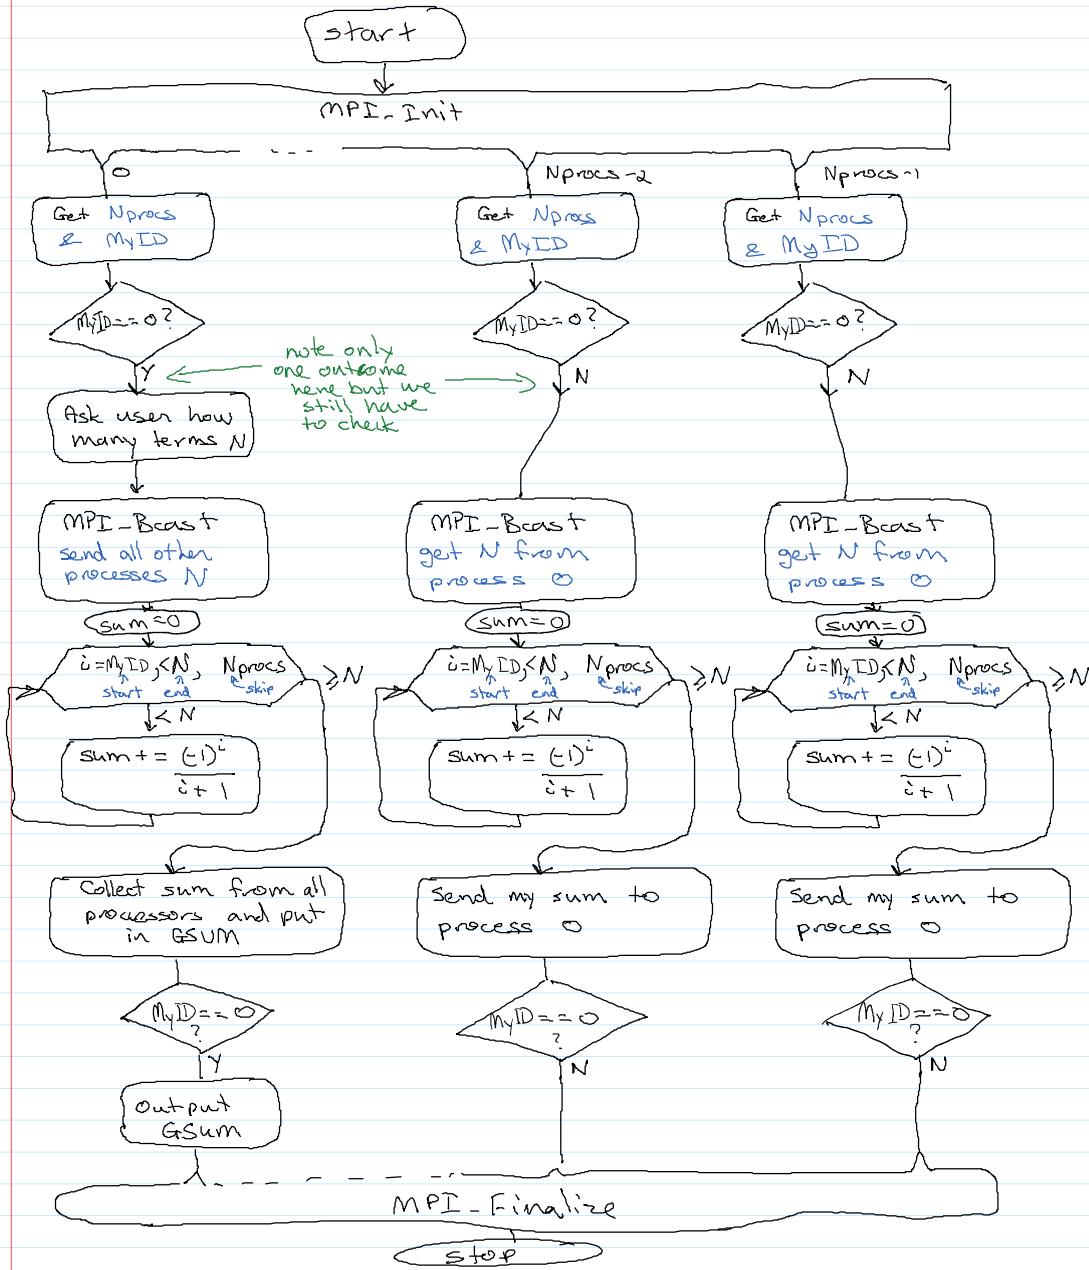
    return 0;
}

```

Same class introduced for helloworld code last lecture

sum different terms based on My ID

And here is a flowchart of what happens in the code



The new commands that take care of the communication are:

`MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD)`

pointer to beginning of data buffer to send
of this size to send (i.e. buffer size)
sending processor who to send to (everyone here)

data size being sent is integer

`MPI_Reduce(&sum, &Gsum, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD)`

data being collected location where to put result
number of this size in buffer operation to perform receiving processor where info coming from

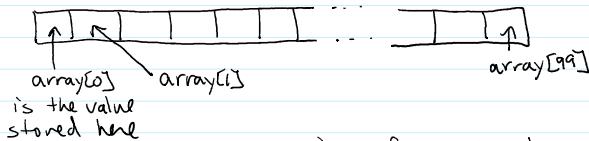
Note: N does not have to be a multiple of Nprocs so some processors will have slightly less work.

So far we have just passed data of a simple fundamental type. To pass larger blocks of data we first need to consider how they are stored.

C-style 1D arrays

double array[100]; creates a C-style array (also in C++)

This is stored in consecutive spaces in memory which we can illustrate as



The size (in memory) of each box is
sizeof(double)

which is 64 bits.

In order to send an array from one process to another we need to know where it starts in memory (its address) and its size.

To get the address of the start of the array we can use the address operator & or the array name

eg

```
#include <iostream>

int main()
{
    int array[5] = { 9, 7, 5, 3, 1 };

    // print the value of the array variable
    std::cout << "The array has address: " << array << '\n';

    // print address of the array elements
    std::cout << "Element 0 has address: " << &array[0] << '\n';

    return 0;
}
```

will produce something like

The array has address: 0042FD5C ← array ↴ note, same address
Element 0 has address: 0042FD5C ← &array[0]

To ensure compatible sizes across a potentially heterogeneous environment, MPI introduces its own fundamental data size → MPI_INT, MPI_FLOAT, MPI_DOUBLE ← use in MPI functions

For dynamic allocation of arrays in C we can use pointers and malloc/calloc. The second is preferable as it also initializes the elements (to zero).

```
float *a;
a = (float *) calloc(n, sizeof(float));
```

data type of a number of elements in array size of each element (in memory)

As with a statically declared array, this creates an array (of size n , which is an integer) stored in a contiguous block of memory. Its elements can be accessed via $a[0]$, $a[1]$, etc as before.

Note: In C (and C++) there is no bounds checking on arrays of this type. So if you try to access $a[i]$ in the first example, or $a[n+2]$ in the second, the compiler will NOT notice. This is a common, and often difficult to spot, source of errors. We will discuss alternatives in later lectures which avoid this, but for now we just need to be careful.

In C++ we can create dynamically allocated arrays using `new[]` and `delete []` but this is no longer recommended (will be made obsolete in future versions of C++). The recommended C++ option is `std::vector`

e.g.

```
std::vector<float> a;
a.reserve(n); a.resize(n);
```

↑ set capacity ↑ set size

will create an array similar to what we did in C. In particular, the elements $a[0], a[1], \dots$ are stored consecutively in memory, which is important for MPI (and performance in general).

A vector has a capacity (total amount of space it might use) and a size (amount of space it is currently using). These can be set (and changed) using the `reserve` and `resize` functions (for capacity and size respectively), with $\text{size} \leq \text{capacity}$.

For scientific computing applications we are often working at the machine limits or need to know the space our code is going to need to run (required by queuing systems on compute canada systems). As such, it is best to manage the capacity manually (e.g. set it when you declare the vector, as in the example above). Also, do NOT use `push-back` to fill the vector).

vector comes with bounds checking, but only if you use $a.at(j)$ rather than $a[j]$. There is a performance hit (small) for using the `.at` function so we don't typically use it. Again, we will introduce a better option later.

One-to-One Communication

e.g. Dot Product

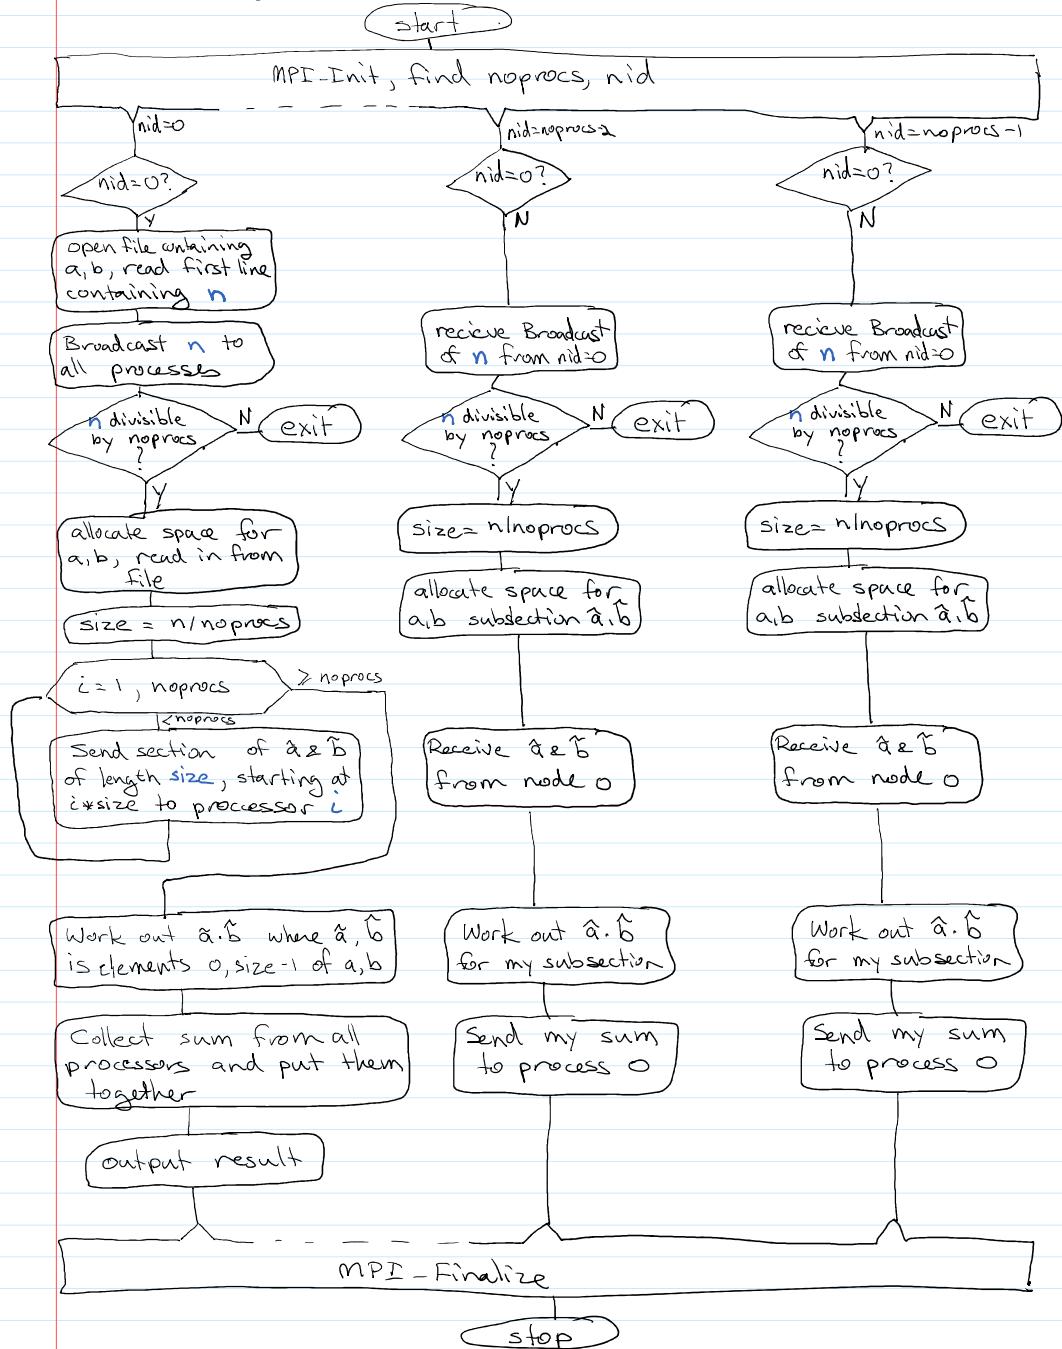
Compute dot product of two vectors "a" and "b" of length n

$$\vec{a} \cdot \vec{b} = \sum_{i=1}^n a_i b_i = \sum_{i=0}^{n-1} a[i] * b[i]$$

let's assume we have N_{procs} processors and that n is exactly divisible by N_{procs} . We can then break up the task into parts of size $= \frac{n}{N_{\text{procs}}}$ as a SIMD problem:

$$\vec{a} \cdot \vec{b} = \sum_{i=0}^{size-1} a[i] * b[i] + \sum_{i=size}^{2*size-1} a[i] * b[i] + \dots + \sum_{i=(N_{\text{procs}}-1)*size}^{n-1} a[i] * b[i]$$

We will again use the Boss-Worker model



```

#include <iostream> // regular I/O
#include <fstream> // file stream header
#include <vector> // needed for std::vector
#include "mpi.h"
  
```

```

class MPI_stuff
{
public:
    int NProcs;
    int MyID;

    MPI_stuff(int argc, char** argv)
    {
        MPI_Init(&argc, &argv);
        MPI_Comm_size(MPI_COMM_WORLD, &NProcs);
    }
}
  
```

MPI class from

```

    {
        MPI_Init(&argc, &argv);
        MPI_Comm_size(MPI_COMM_WORLD, &NProcs);
        MPI_Comm_rank(MPI_COMM_WORLD, &MyID);
    }

    ~MPI_stuff()
    {
        MPI_Finalize();
    }
}

```

MPI class from before

```

// Get the number of rows of data in the file from first row of file
int GetNumberElements(std::ifstream &fin)
{
    int n;

    if (fin.is_open())
        fin >> n;
    else { // no file to read from
        std::cout << "Input File not found\n";
        MPI_Abort(MPI_COMM_WORLD,-1);
    }
    return n;
}

```

function to get n from first entry of a file

```

// Read n elements of two column data from a file, store in arrays a,b
void ReadArrays(std::ifstream &fin, std::vector<float> &a, std::vector<float> &b, int n)
{
    if(fin.is_open()){
        for(int i = 0; i < n; i++)
            fin >> a[i] >> b[i];
    }
    else{ // fp null means no file to read from
        std::cout << "Input File not found\n";
        MPI_Abort(MPI_COMM_WORLD,-1);
    }
}

```

read arrays from a file

```

// Dot product of two vectors owned and fully stored on Boss node
float DotProduct(std::vector<float> &a, std::vector<float> &b,
                 int size, const MPI_stuff &the_mpi)
{
    MPI_Status status;

```

BOSS sends parts of a,b to workers

matching calls

```

    // Distribute or Scatter the array to the workers
    if(the_mpi.MyID == 0){
        for(int i = 1; i < the_mpi.NProcs; i++){
            MPI_Send(&a[size*i],size,MPI_FLOAT,i,10,MPI_COMM_WORLD);
            MPI_Send(&b[size*i],size,MPI_FLOAT,i,20,MPI_COMM_WORLD);
        }
    } else{
        MPI_Recv(&a[0],size,MPI_FLOAT,0,10,MPI_COMM_WORLD,&status);
        MPI_Recv(&b[0],size,MPI_FLOAT,0,20,MPI_COMM_WORLD,&status);
    }
}

```

actual a·b done here

```

    // Work out my part of the sum
    float sum=0, Gsum=0;
    for(int i = 0; i < size; i++)
        sum += a[i] * b[i];
}

```

} work out dot product of vector fragments

```

    // Collect results in Gsum
    MPI_Reduce(&sum,&Gsum,1,MPI_FLOAT,MPI_SUM,0,MPI_COMM_WORLD); < collect and sum results

    return Gsum;
}

```

```

/* See lecture notes for comments */
int main(int argc, char** argv)
{
    MPI_stuff the_mpi(argc, argv);

    std::ifstream fin;
    int n = 0;

    if(the_mpi.MyID == 0) {
        fin.open("DotData.txt");
        n = GetNumberElements(fin);
    }
    MPI_Bcast(&n, 1, MPI_INT, 0, MPI_COMM_WORLD); ← send n to everyone
    } only Boss node reads from the file
    } (multiple processes cannot access file simultaneously)

    if(n % the_mpi.NProcs) {
        std::cout << "Number of processes is not a multiple of n.\n";
        MPI_Abort(MPI_COMM_WORLD, -2); ← kills program gracefully
    } ← is n divisible by NProcs,
    } } IF not exit program
    } on all nodes

    int size = n / the_mpi.NProcs;

    std::vector<float> a, b;
    if(the_mpi.MyID == 0) {
        a.reserve(n); // try to avoid std::vector reallocations } node 0 stores full
        b.reserve(n); arrays
        ReadArrays(fin, a, b, n); ← node 0 reads from file
    }
    else {
        a.reserve(size); a.resize(size);
        b.reserve(size); b.resize(size);
    } } worker nodes only need an array
    } that fits the part they are working on

    float adotb = DotProduct(a, b, size, the_mpi); ← compute a · b

    if(the_mpi.MyID == 0)
        std::cout << "The inner product is " << adotb << std::endl; ← node 0 (Boss) does output
    } of result
    return 0;
}

```

This program introduces some new MPI commands:

pointer to beginning
of buffer to be sent

processor to send to

"MPI communicator"

`MPI_Send(&a[size*i],size,MPI_FLOAT,i,10,MPI_COMM_WORLD);`

number of size in buffer

tag for this message (to match Recv)

these need to match (b given different tag)

pointer to beginning of receive buffer

where to get message from

tag

`MPI_Recv(&a[0],size,MPI_FLOAT,0,10,MPI_COMM_WORLD,&status);`

Comments

1. We could fix the $n \% nprocs$ condition several ways.
An easy fix would be zero padding the end of the array.
2. Boss node uses more memory than other nodes. We could fix this by reading segments of the file and then sending before reading next segment.
If at the same time you got the Boss to work on the end of a[b] (rather than beginning)
you could solve problem ① by having Boss do the

- you could solve problem ① by having Boss do the "special" array of smaller size at the end.
3. Potential crash points should be noted and have program exit gracefully with MPI_Abort. This will ensure all MPI processes are stopped.

Summary of MPI Commands encountered so far

```

MPI_Init(&argc, &argv);
MPI_Finalize();
MPI_Abort(MPI_COMM_WORLD, -1);

MPI_Comm_size(MPI_COMM_WORLD, &nprocs);
MPI_Comm_rank(MPI_COMM_WORLD, &myrank);

MPI_Bcast(&N, 1, MPI_INT, 0, MPI_COMM_WORLD);
MPI_Reduce(&sum, &Gsum, 1, MPI_FLOAT, MPI_SUM, 0, MPI_COMM_WORLD);

MPI_Send(&a[size*i], size, MPI_FLOAT, i, 10, MPI_COMM_WORLD);
MPI_Recv(&a[0], size, MPI_FLOAT, 0, 10, MPI_COMM_WORLD, &status);

```

Another variant of MPI_Reduce is MPI_Allreduce, e.g

```
MPI_Allreduce(&maxerr, &maxerrG, 1, MPI_FLOAT, MPI_MAX, MPI_COMM_WORLD);
```

In this case maxerr gets reduced to All processors, not only to one specific node (result in "maxerrG"). The operations (of type MPI_OP) are varied, in this case it finds the maximum of all maxerr and puts it into maxerrG. (This is All-to-All communication).

The following predefined operations are supplied for MPI_Reduce and related functions [MPI_Allreduce](#), [MPI_Reduce_scatter](#), and [MPI_Scan](#). These operations are invoked by placing the following in op:

Name	Meaning
MPI_MAX	maximum
MPI_MIN	minimum
MPI_SUM	sum
MPI_PROD	product
MPI_LAND	logical and
MPI_BAND	bit-wise and
MPI_LOR	logical or
MPI_BOR	bit-wise or
MPI_LXOR	logical xor
MPI_BXOR	bit-wise xor
MPI_MAXLOC	max value and location
MPI_MINLOC	min value and location

The two operations MPI_MINLOC and MPI_MAXLOC are discussed separately below (MINLOC and MAXLOC).

From https://www.open-mpi.org/doc/v1.8/man3/MPI_Reduce.3.php

See the source above for more (OpenMPI docs).

Before going on let's discuss the "status" variable in MPI_Recv

MPI_Status is a data structure type with fields:

MPI_SOURCE - id of processor sending message
 MPI_TAG - the message tag

MPI_Status is a data structure type with fields:

- MPI_SOURCE - id of processor sending message
- MPI_TAG - the message tag
- MPI_ERROR - error status (often not modified as would be the same as what function returns)

other fields are possible but implementation dependent
 So `status.MPI_SOURCE` above could be used to get sender id.

Controlling Synchronization - Blocking & Non Blocking Communication

MPI_Send is what is known as a "blocking" routine:
 The sending processor waits for a "received" confirmation before working on the next line of code. Similarly
MPI_Recv causes the receiving processor to wait for the message before moving on.

e.g. Suppose I had 10 tasks, numbered from 1 to 10,

to distribute to $nprocs = 4$ so $node \in \{0, 1, 2, 3\}$

all integers

```
for (task=1; task <= 10; task++) {
    node=task-nprocs*(task/nprocs); ← everything here is an integer
    ...
}
```

so this expression is done using integer arithmetic

task	<u>(task / nprocs)</u>	node
1	0	1
2	0	2
3	0	3
4	1	0
5	1	1
6	1	2
7	1	3
8	2	0
9	2	1
10	2	2

So now consider the following code:

```
#include <stdio.h>
#include "mpi.h"
...

```

```

main(int argc, char** argv)
{
    struct mpi_vars this_mpi = mpi_start(argc, argv);
    MPI_Status status;
    int task, taskx, node;

    for (task=1; task <= 10; task++) {
        node=task-this_mpi.NProcs*(task>this_mpi.NProcs); see which node should do the task

        if (node == this_mpi.MyID && node != 0) { //Check to see if this is my task to do
            if (node==2) waste_your_time(); some function that takes up some time
            MPI_Send(&task,1,MPI_INT,0,10,MPI_COMM_WORLD); //Tell Boss I am the one
        }
        if (this_mpi.MyID == 0) {
            if (node == 0)
                printf("Processor %d will compute %d\n", 0, task);
            else {
                MPI_Recv(&taskx,1,MPI_INT,MPI_ANY_SOURCE,10,MPI_COMM_WORLD,&status); take a message from anyone
                printf("Processor %d will compute %d\n", status.MPI_SOURCE, taskx); node the message actually came from
            }
        }
    }

    for (task=11; task <= 20; task++) {
        node=task-this_mpi.NProcs*(task>this_mpi.NProcs);

        if (node == this_mpi.MyID && node != 0) { //Check to see if this is my task to do
            if (node==2) waste_your_time();
            MPI_Send(&task,1,MPI_INT,0,10,MPI_COMM_WORLD); //Tell Boss I am the one
        }
        if (this_mpi.MyID == 0) {
            if (node == 0)
                printf("Processor %d will compute %d\n", 0, task);
            else {
                MPI_Recv(&taskx,1,MPI_INT,MPI_ANY_SOURCE,10,MPI_COMM_WORLD,&status);
                printf("Processor %d will compute %d\n", status.MPI_SOURCE, taskx);
            }
        }
    }
    MPI_Finalize();
}

```

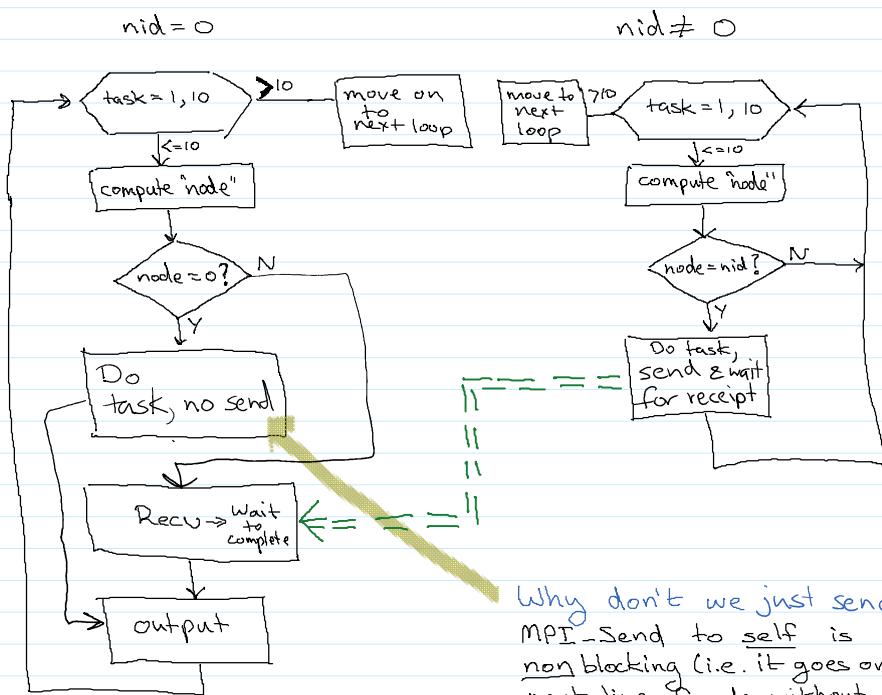
The output from such a code looks something like:

Processor	1 will compute	1
Processor	1 will compute	5
Processor	1 will compute	9
Processor	1 will compute	13
Processor	1 will compute	17
Processor	3 will compute	3
Processor	3 will compute	7
Processor	3 will compute	11
Processor	3 will compute	15
Processor	3 will compute	19
Processor	0 will compute	4
Processor	0 will compute	8
Processor	0 will compute	12
Processor	2 will compute	2
Processor	2 will compute	6
Processor	0 will compute	16
Processor	2 will compute	10
Processor	2 will compute	14
Processor	2 will compute	18
Processor	0 will compute	20

Comments:

1. The tasks are not completed in order (and order may change if you rerun it)
2. Node 2 is slower to complete tasks (as expected from `waste_your_time()` call.)
3. node 0 is slower to complete tasks because it is busy getting results from other processors and outputting results.
4. node 1 & 3 are finished all tasks before node 0 & 2 are done their first task (they complete both loops before 0&2 finish 1st loop).

Let's create an old-style flowchart for this code's 1st loop:



Why don't we just send to self here?

MPI-Send to self is non blocking (i.e. it goes on with next line of code without waiting for a "received" confirmation.)

However, some implementations have a bug on this and block.

If we need to have all of tasks 1-10 complete before going on to the 11-20 tasks, we need to synchronize the threads by adding an MPI-Barrier line between the loops:

`MPI_Barrier (MPI_COMM_WORLD);`

- blocks caller from proceeding until all processes in MPI_COMM_WORLD have made a similar call.

Adding this line between the loops results in the following output:

Processor	1 will compute	1
Processor	3 will compute	3
Processor	1 will compute	5
Processor	3 will compute	7
Processor	1 will compute	9
Processor	0 will compute	4
Processor	2 will compute	2
Processor	0 will compute	8
Processor	2 will compute	6
Processor	2 will compute	10
Processor	1 will compute	13
Processor	3 will compute	11
Processor	1 will compute	17
Processor	3 will compute	15
Processor	3 will compute	19
Processor	0 will compute	12
Processor	0 will compute	16
Processor	2 will compute	14
Processor	2 will compute	18
Processor	0 will compute	20

As MPI_Send is a blocking function it is possible to force the tasks to be performed in order by making the "tag" more specific. i.e: change loop to:

```
for (task=1; task <= 10; task++) {
    node=task-this_mpi.NProcs*(task/this_mpi.NProcs);

    if (node == this_mpi.MyID && node != 0) { //Check to see if this is my task to do
        if (node==2) waste_your_time();
        MPI_Send(&task,1,MPI_INT,0,10+node,MPI_COMM_WORLD); //Tell Boss I am the one
    }
    if (this_mpi.MyID == 0) {
        if (node == 0)
            printf("Processor %d will compute %d\n", 0, task);
        else {
            MPI_Recv(&taskx,1,MPI_INT,MPI_ANY_SOURCE,10+node,MPI_COMM_WORLD,&status);
            printf("Processor %d will compute %d\n", status.MPI_SOURCE, taskx);
        }
    }
}
```

Note however that this will slow things down as a processor that finishes early will now have to wait for node 0 to receive the previous task before moving on to the next task \Rightarrow effectively no parallelization as tasks are done sequentially in time, just on different processors.

Non-blocking Communication

```
MPI_Isend(&old[size-1][1],N-1,MPI_FLOAT,nid+1,10,MPI_COMM_WORLD,&req_send10)
MPI_Irecv(&old[0][1],N-1,MPI_FLOAT,nid+1,10,MPI_COMM_WORLD,&req_recv10)
MPI_Request
MPI_REQUEST_NULL
MPI_Wait(&req_recv10,&status)
```

An immediate send (MPI_Isend) implies that the sending processor does not have to wait for a "received" confirmation before working on the next line of code. Such confirmation would be required for the regular MPI_Send. MPI_Isend is a nonblocking routine, while MPI_Send is a blocking routine.

In MPI_Isend note in particular the pointer to the variable req_send10. This is an object of the type MPI_Request. This object can be accessed again later to see whether the message has been delivered.

Similarly, MPI_Irecv contains a request handle, namely req_recv10 so that completion can be checked later.

`MPI_REQUEST_NULL` is the “default” value with which any variable/object of the request type should be initialized.

Lastly, `MPI_Wait` makes the executing node wait until the task associated with `req_recv10` is successfully finished.

Send requests often do not need to have confirmation, whereas if you are receiving data you need to check whether you got something before working with the data.
Note that `MPI_IRecv + MPI_Wait` is effectively the same as `MPI_Recv`.

e.g. Changing our loop from previous example to:

```
for (task=1; task <= 10; task++) {  
    node=task-this_mpi.NProcs*(task>this_mpi.NProcs);  
  
    if (node == this_mpi.MyID) { //Check to see if this is my task to do  
        if (node==2) waste_your_time();  
        MPI_Isend(&task,1,MPI_INT,0,10,MPI_COMM_WORLD,&req_send10); //Tell Boss I am the one  
    }  
    if (this_mpi.MyID == 0) {  
        req_recv10 = MPI_REQUEST_NULL;  
        MPI_Irecv(&taskx,1,MPI_INT,MPI_ANY_SOURCE,10,MPI_COMM_WORLD,&req_recv10);  
        MPI_Wait(&req_recv10,&status);  
        printf("Processor %d will compute %d\n", status.MPI_SOURCE, taskx);  
    }  
}
```

need to set how so not prematurely passed

previously declared of type MPI_Request

This removes the blocking effect of `MPI_Send` and allows all processes (including `MyID=0`) to proceed through the entire list of tasks without waiting for BOSS processor to confirm receipt.

To avoid having BOSS node waiting for others to complete before doing its task, we should move the receives to a separate statement:

```
for (task=1; task <= 10; task++) {  
    node=task-this_mpi.NProcs*(task>this_mpi.NProcs);  
  
    if (node == this_mpi.MyID) { //Check to see if this is my task to do  
        if (node==2) waste_your_time();  
        MPI_Isend(&task,1,MPI_INT,0,10,MPI_COMM_WORLD,&req_send10); //Tell Boss I am the one  
    }  
  
    if (this_mpi.MyID == 0) {  
        for (task=1; task <= 10; task++) {  
            req_recv10 = MPI_REQUEST_NULL;  
            MPI_Irecv(&taskx,1,MPI_INT,MPI_ANY_SOURCE,10,MPI_COMM_WORLD,&req_recv10);  
            MPI_Wait(&req_recv10,&status);  
            printf("Processor %d will compute %d\n", status.MPI_SOURCE, taskx);  
        }  
    }  
}  
MPI_Barrier(MPI_COMM_WORLD); a good idea to put this here to prevent more messages being sent to nid=0 by other processors and getting confused with these tasks by nid=0 as the IRecv is very open (i.e. non-specific, it is taking messages from anyone).
```